



# 3장 *Getting started with neural networks*

“기회와 준비가 만났을 때 ... ”



# Outline



- Core components of **neural networks**
- An introduction to **Keras**
- Setting up a deep-learning workstation
- Using neural networks to solve basic **classification** and **regression** problems



# Outline



- Classifying movie reviews as positive or negative (binary classification) - **IMDB**



- Classifying news wires by topic (multiclass classification) - **Reuters**



- Estimating the price of a house, given real-estate data (regression) – **Boston Housing**



# 1. Anatomy of a neural network

- Relationship between the network, **layers**, **loss function**, and **optimizer**

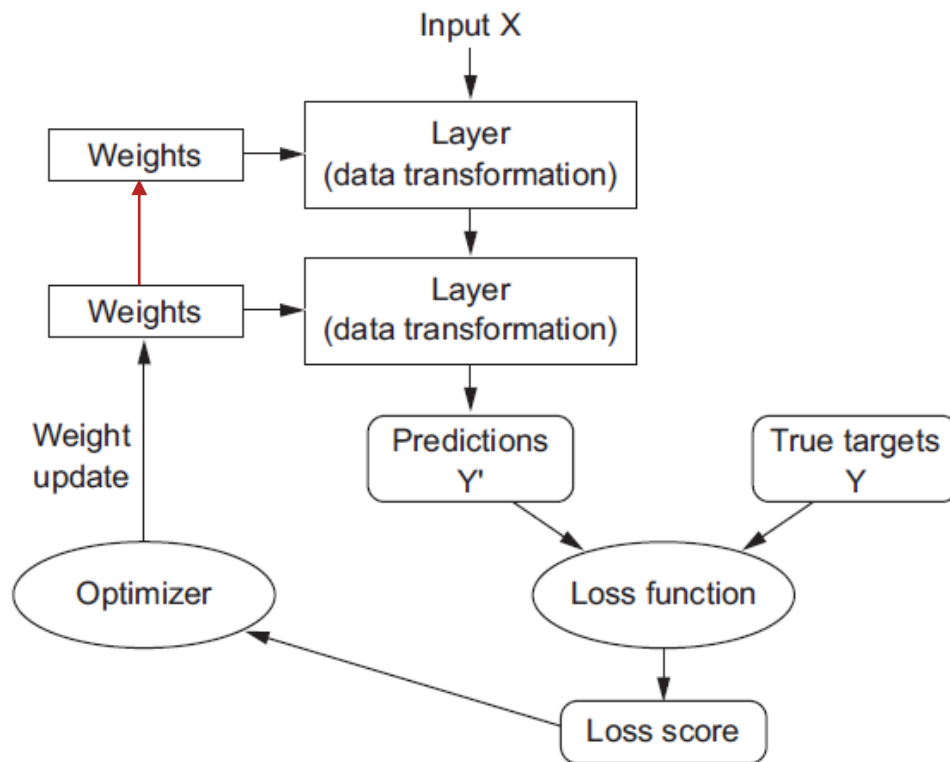


Figure 3.1 Relationship between the network, layers, loss function, and optimizer

# ○○○ 1. Anatomy of a neural network ○○○

## 3.1.1 *Layers*: the building blocks of deep learning

- a data-processing module
- input and outputs tensors
- *weights* tensors contain the network's *knowledge* - 학습결과
- densely connected layers, fully connected, or dense layers (the **Dense** class in Keras)

# 1. Anatomy of a neural network

## 3.1.1 Layers: the building blocks of deep learning

# A dense layer with 10 **output units**

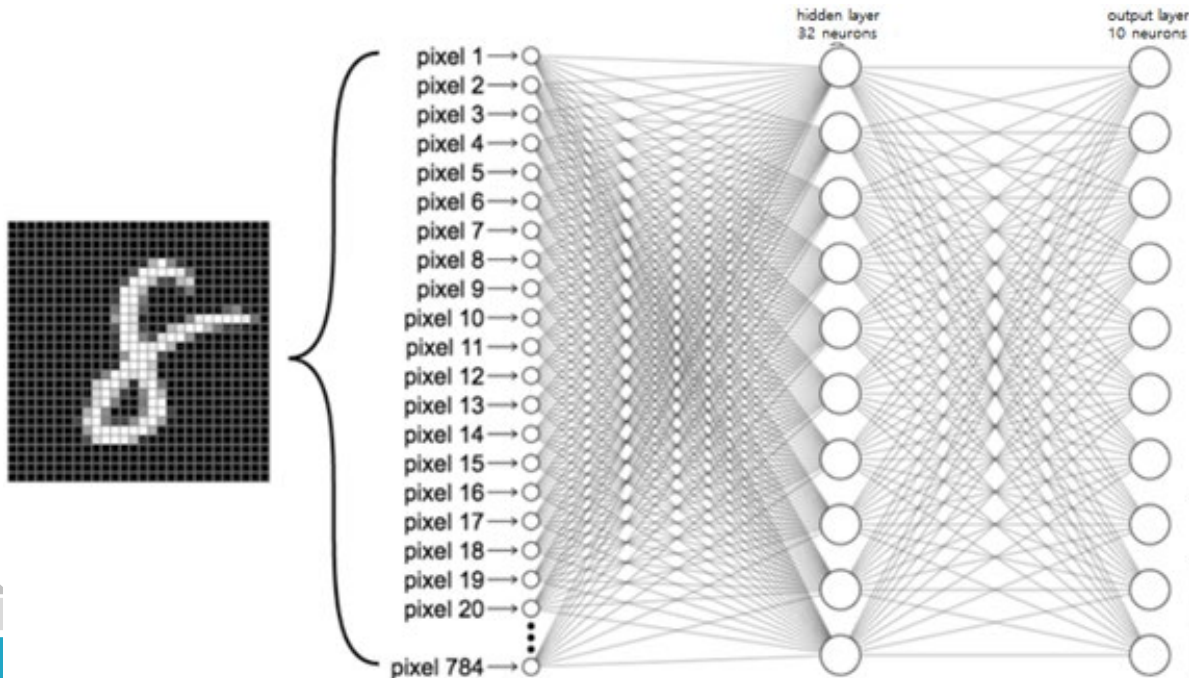
```
from keras import models
```

```
from keras import layers
```

```
model = models.Sequential()
```

```
model.add(layers.Dense(32, input_shape=(784,)))
```



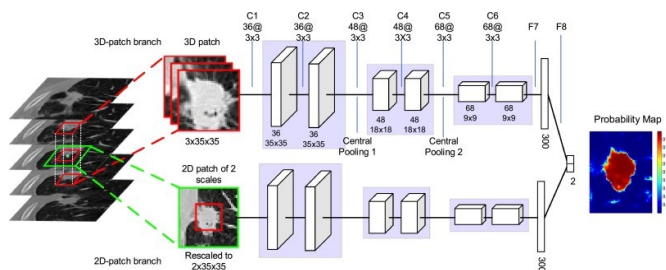
```
model.add(layers.Dense(10))
```



# 1. Anatomy of a neural network

## 3.1.2 Models: networks of layers

- ▶ A deep-learning model is a **directed, acyclic graph** of layers.
- ▶ Search a good set of values for the **weight tensors**
- ▶ Some common **topologies** include the following:
  - Two-branch networks
  - Multihead networks
  - Inception blocks - Concatenate



# ○○○ 1. Anatomy of a neural network ○○○

## 3.1.3 *Loss functions and optimizers:* keys to configuring the *learning process*

- ▶ choose two more things:
  - *Loss function* (*objective function*)—The quantity that will be minimized during training. It represents a measure of success for the task at hand.
  - *Optimizer*—Determines how the network will be updated based on the loss function. It implements a specific variant of stochastic gradient descent (SGD).



## 2. Introduction to Keras

- ▶ code examples use Keras (<https://keras.io>)
- ▶ Keras is a **deep-learning framework for Python** that provides a convenient way to define and train almost any kind of deep-learning model.
- ▶ Keras was initially developed for researchers, with the aim of enabling **fast experimentation**.
- ▶ Keras has the following key features:
  - **same code** on CPU or GPU.
  - It has a **user-friendly API** that makes it easy to quickly prototype deep-learning models.
  - It has **built-in support** for convolutional networks, recurrent networks
  - It supports arbitrary network architectures: multi-input or multi-output models, layer sharing, model sharing, and so on.

## 2. Introduction to Keras

- ▶ Keras is distributed under the permissive **MIT license**, which means it can be **freely used in commercial projects**.
- ▶ Keras is used at **Google, Netflix, Uber, CERN, Yelp(클라우드 소싱 리뷰 포럼), Square(모바일 결제 기업)**, and hundreds of startups working on a wide range of problems.
- ▶ Keras is also a popular **framework on Kaggle**, the machine-learning competition website, where almost every recent deep-learning competition has been won using Keras models.


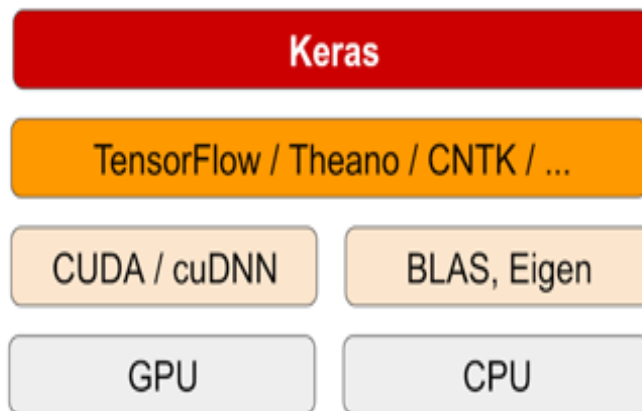


Figure 3.2 Google web search interest for different deep-learning frameworks over time

## 2. Introduction to Keras

### 3.2.1 Keras, TensorFlow, Theano, and CNTK

- ▶ Keras is a **model-level library**, providing high-level building blocks for developing deep-learning models.
- ▶ the three existing backend implementations - the **TensorFlow**, the **Theano** (Quebec AI Institute, University of Montreal), and the Microsoft Cognitive Toolkit (**CNTK**) backends. In the future, it's likely that Keras will be extended to work with even more deep-learning execution engines.
- ▶ **Caffe** - Berkeley Vision and Learning Center



**Figure 3.3** The deep-learning software and hardware stack

## 2. Introduction to Keras

### 3.2.1 Developing with Keras: a quick overview

- ▶ The typical Keras workflow looks just like that example:
  - 1 Define your **training data**: **input tensors** and **target tensors**.
  - 2 Define **a network** of **layers** (or *model*) that maps your inputs to your targets
  - 3 Configure the learning process by choosing a **loss function**, an **optimizer**, and some **metrics** to monitor (`compile()`).
  - 4 Iterate on your training data by calling the `fit()` method of your model.
- ▶ There are two ways to define a model: `Sequential` class, the *functional API* (for **directed acyclic graphs** of layers, which lets you build completely **arbitrary architectures**).

## 2. Introduction to Keras

### 3.2.1 Developing with Keras: a quick overview

► a two-layer model defined using the `Sequential` class :

```
from keras import models
from keras import layers

model = models.Sequential()
model.add(layers.Dense(32, activation='relu',
                       input_shape=(784,)))
model.add(layers.Dense(10, activation='softmax'))
```

## 2. Introduction to Keras

### 3.2.1 Developing with Keras: a quick overview

- ▶ in the compilation step, where you specify the optimizer and loss function(s)

```
from keras import optimizers  
model.compile(optimizer=optimizers.RMSprop(lr=0.001),  
loss='mse', metrics=['accuracy']) # lr →  $\eta$  (Eta)
```

- ▶ the `fit()` method:

```
model.fit(input_tensor, target_tensor, batch_size=128, epochs=10)
```

## 4. Classifying movie reviews: a binary classification example

### 3.4 Classifying movie reviews: a binary classification example

- ▶ Two-class classification, or binary classification - classify movie reviews as positive or negative, based on the text content of the reviews

#### 3.4.1 The IMDB dataset

- ▶ IMDB dataset: a set of 50,000 highly polarized reviews from the *Internet Movie Database*.
- ▶ They're split into 25,000 reviews for training and 25,000 reviews for testing, each set consisting of 50% negative and 50% positive reviews.
- ▶ IMDB dataset: the reviews (sequences of words) into sequences of integers

## 4. Classifying movie reviews: a binary classification example

### 3.4.1 The IMDB dataset

- ▶ The following code will load the dataset (when you run it the first time, about 80 MB of data will be downloaded to your machine).

# of training data – 25,000

# of test data – 25,000

#### Listing 3.1 Loading the IMDB dataset

```
from keras.datasets import imdb

(train_data, train_labels), (test_data, test_labels)
    = imdb.load_data(num_words=10000)

# top 10,000 most frequently occurring words in the training data
>>> train_data[0]
[1, 14, 22, 16, ... 178, 32] # list of word indices
>>> train_labels[0]
1 # 0 stands for negative and 1 stands for positive
```



## 4. Classifying movie reviews: a binary classification example

### 3.4.1 The IMDB dataset

- ▶ no **word index** will exceed 10,000:

```
>>> max([max(sequence) for sequence in  
train_data])  
9999
```

- ▶ back to English words:

```
word_index = imdb.get_word_index()  
reverse_word_index = dict(  
    [(value, key) for (key, value) in word_index.items()])
```

**dict**

```
reverse_word_index = dict([(value, key) for (key, value) in word_index.items()])  
print(sorted(reverse_word_index.items()))
```

```
[(1, 'the'), (2, 'and'), (3, 'a'), (4, 'of'), (5, 'to'), (6, 'is'), (7, 'br'), (8,  
'in'), (9, 'it'), (10, 'i'), (11, 'this'), (12, 'that'), (13, 'was'), (14, 'as'),  
(15, 'for'), (16, 'with'), (17, 'movie'), (18, 'but'), (19, 'film'), (20, 'on'),  
(21, 'by'), (22, 'he'), (23, 'she'), (24, 'it'), (25, 'we'), (26, 'you'), (27,
```

## 4. Classifying movie reviews: a binary classification example

### 3.4.1 The IMDB dataset

▶ no **word index** will exceed 10,000:

```
decoded_review = ' '.join(
    [reverse_word_index.get(i - 3, '?') for i in train_data[0]])
# 0, 1, 2는 '패딩', '문서 시작', '사전에 없음' 을 위한 인덱스이므로 3을 뺍니다.
```

```
print(train_data[11])
```

```
[1, 54, 13, 1610, 14, 20, 13, 69, 55, 364, 1398, 21, 54, 13, 219, 12, 13, 1706, 15, 4, 20, 16, 329, 6, 176, 329, 74,
51, 13, 873, 4, 156, 71, 78, 4, 7412, 322, 16, 31, 7, 4, 249, 4, 65, 16, 38, 379, 12, 100, 157, 18, 6, 910, 20, 549,
18, 4, 1496, 21, 14, 31, 9, 24, 6, 212, 12, 9, 6, 1322, 991, 7, 3002, 4, 425, 9, 73, 2218, 549, 18, 31, 155, 36, 100,
763, 379, 20, 103, 351, 5308, 13, 202, 12, 2241, 5, 6, 320, 46, 7, 457]
```

? when i rented this movie i had very low expectations but when i saw it i realized that the movie was less a lot less than what i expected the actors were bad the doctor's wife was one of the worst the story was so stupid it could work for a disney movie except for the murders but this one is not a comedy it is a laughable masterpiece of stupidity the title is well chosen except for one thing they could add stupid movie after dead husbands i give it 0 and a half out of 5

## 4. Classifying movie reviews: a binary classification example

### 3.4.2 Preparing the data

- ▶ turn your lists into tensors - **vectorize the data**

[1, 14, 22, 16, ... 178, 32] **review words** →

[0., 1., 0., ..., 1., ..., 0., 0.] 신경망의 입력을 위한 일정한 10,000개 원소로 된 벡터

#### Listing 3.2 Encoding the integer sequences into a binary matrix

```
import numpy as np
```

```
def vectorize_sequences(sequences, dimension=10000):  
    results = np.zeros((len(sequences), dimension))  
    for i, sequence in enumerate(sequences):  
        results[i,] = 1. # [... , 3, ...] → [0., 0., 0., 1.,...]  
    return results
```

```
x_train = vectorize_sequences(train_data)
```

```
x_test = vectorize_sequences(test_data)
```


```
>>> x_train[0]  
array([ 0.,  1.,  1., ...,  0.,  0.,  0.])
```

- ▶ **vectorize labels:**

```
y_train = np.asarray(train_labels).astype('float32')
```

```
y_test = np.asarray(test_labels).astype('float32')
```

Vocabulary:  
Man, woman, boy,  
girl, prince,  
princess, queen,  
king, monarch



	1	2	3	4	5	6	7	8	9
man	1	0	0	0	0	0	0	0	0
woman	0	1	0	0	0	0	0	0	0
boy	0	0	1	0	0	0	0	0	0
girl	0	0	0	1	0	0	0	0	0
prince	0	0	0	0	1	0	0	0	0
princess	0	0	0	0	0	1	0	0	0
queen	0	0	0	0	0	0	1	0	0
king	0	0	0	0	0	0	0	1	0
monarch	0	0	0	0	0	0	0	0	1

Each word gets  
a 1x9 vector  
representation

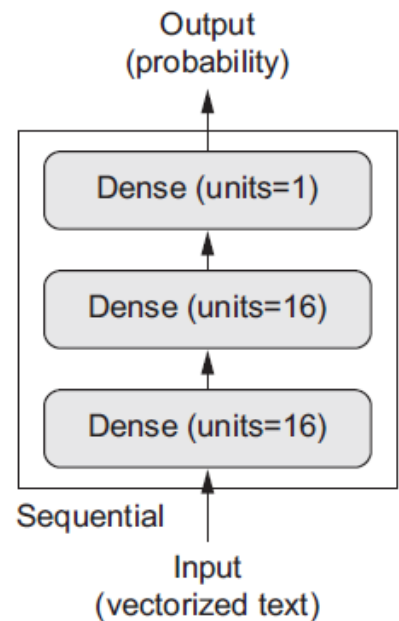


## 4. Classifying movie reviews: a binary classification example

### 3.4.3 Building your network

#### ► Listing 3.3 The model definition

```
from keras import models from keras import layers

model = models.Sequential()
model.add(layers.Dense(16, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(16, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid')) #binary classification
```




## 4. Classifying movie reviews: a binary classification example

### 3.4.3 Building your network

#### Listing 3.4 Compiling the model

```
model.compile(optimizer='rmsprop',  
              loss='binary_crossentropy',  
              metrics=['accuracy'])
```



## 4. Classifying movie reviews: a binary classification example

### 3.4.3 Building your network

#### Listing 3.5

##### ► Configuring the optimizer

```
from keras import optimizers

model.compile(optimizer=optimizers.RMSprop(lr=0.001),
              loss='binary_crossentropy', metrics=['accuracy'])
```

#### Listing 3.6 Using custom losses and metrics

```
from keras import losses
from keras import metrics

model.compile(optimizer=optimizers.RMSprop(lr=0.001),
              loss=losses.binary_crossentropy,
              metrics=[metrics.binary_accuracy])
```

## 4. Classifying movie reviews: a binary classification example

### 3.4.5 Validating your approach

- ▶ create a validation set by setting apart 10,000 samples from the original training data.

#### Listing 3.7 Setting aside a validation set

```
x_val = x_train[:10000] # 검증 데이터
partial_x_train = x_train[10000:] # 훈련 데이터
y_val = y_train[:10000] # 검증 label
partial_y_train = y_train[10000:] # 훈련 label
```

#### Listing 3.8 Training your model

```
model.compile(optimizer='rmsprop',
              loss='binary_crossentropy', metrics=['acc'])
history = model.fit(partial_x_train,
                    partial_y_train, epochs=20, batch_size=512,
                    validation_data=(x_val, y_val))
```

```
>>> history_dict = history.history # model.fit 훈련정보를 dictionary에 반환
>>> history_dict.keys()
[u'acc', u'loss', u'val_acc', u'val_loss']

>>> acc = history_dict['acc']
```

## 4. Classifying movie reviews: a binary classification example

### 3.4.4 Validating your approach


► **Matplotlib** to plot the training and validation **loss** side by side (see figure 3.7)

#### Listing 3.9 Plotting the training and validation loss

```
import matplotlib.pyplot as plt
history_dict = history.history
loss_values = history_dict['loss']
val_loss_values = history_dict['val_loss']
epochs = range(1, len(acc) + 1)

plt.plot(epochs, loss_values, 'bo', label='Training loss')
plt.plot(epochs, val_loss_values, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```

Figure 3.7 Training and validation **loss**





## 4. Classifying movie reviews: a binary classification example

### 3.4.4 Validating your approach

► **Matplotlib** to plot the training and validation **accuracy** (see figure 3.8)

#### Listing 3.10 Plotting the training and validation accuracy

```
plt.clf() # 생성한 그래프를 clear
acc = history_dict['acc']
val_acc = history_dict['val_acc']

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Acc')
plt.legend()
plt.show()
```




Figure 3.8 Training and validation **accuracy**

## 4. Classifying movie reviews: a binary classification example

### 3.4.4 Validating your approach

► *overfitting*: specific to the training data and don't generalize to data outside of the training set

#### Listing 3.11 Retraining a model from scratch

```
model = models.Sequential()
model.add(layers.Dense(16, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(16, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))

model.compile(optimizer='rmsprop',
              loss='binary_crossentropy', metrics=['accuracy'])

model.fit(x_train, y_train, epochs=4, batch_size=512)
results = model.evaluate(x_test, y_test)
```

► The final results are as follows:

```
>>> results
[0.2929924130630493, 0.8832799999999999]
```

## 4. Classifying movie reviews: a binary classification example

### 3.4.5 Using a trained network to generate predictions on new data

- ▶ After having trained a network, you'll want to use it in a practical setting. You can generate the likelihood of reviews being positive by using the **predict** method:

```
>>> model.predict(x_test) # 25,000
array([[ 0.98006207]
       [ 0.99758697]
       [ 0.99975556]
       ...,
       [ 0.82167041]
       [ 0.02885115]
       [ 0.65371346]], dtype=float32)
```

As you can see, the network is confident for some samples (0.99 or more, or 0.01 or less) but less confident for others (0.6, 0.4).

## 4. Classifying movie reviews: a binary classification example

### 3.4.7 Wrapping up

- **preprocessing** on your raw data in order to be able to feed it—as tensors—into a neural network. Sequences of words can be encoded as **binary vectors**.
- Stacks of **Dense** layers with **relu** activations can solve a wide range of problems (including sentiment classification).
- In a **binary classification problem** (two output classes) - end with **one unit Dense** layer and a **sigmoid activation**: the **output** of your network should be a scalar **between 0 and 1**, encoding a probability.
- a **binary classification problem** - use **binary\_crossentropy** loss function
- The **rmsprop optimizer** is generally a good enough choice
- **overfitting** - worse results on data they've **never seen before**. Be sure to always monitor performance on data that is **outside of the training set**.



## 2. Introduction to Keras



- 추가 실험 예
  - 레이어의 수
  - 레이어의 node 수
  - 다양한 활성화 함수
  - 다양한 loss 함수
  - `num_words=5000`
  - `model.fit(x_train, y_train, epochs=10, batch_size=100)`
  - `model.summary()` 추가 결과