


5장 *Deep learning for computer vision*

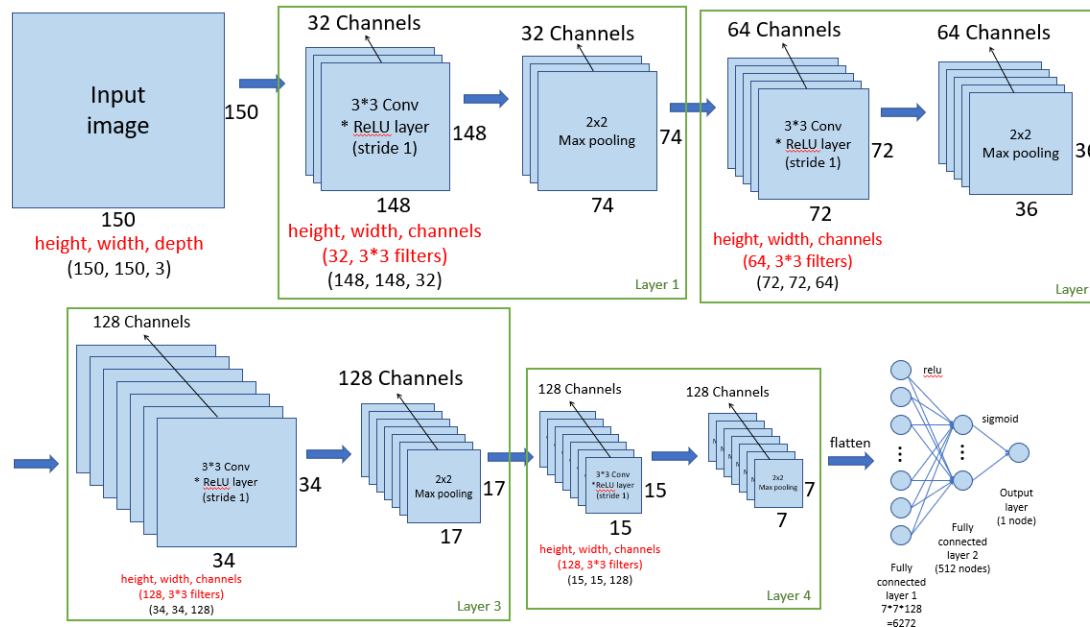
“Out of the Box”



5.2 Training a convnet from scratch on a small dataset

실습

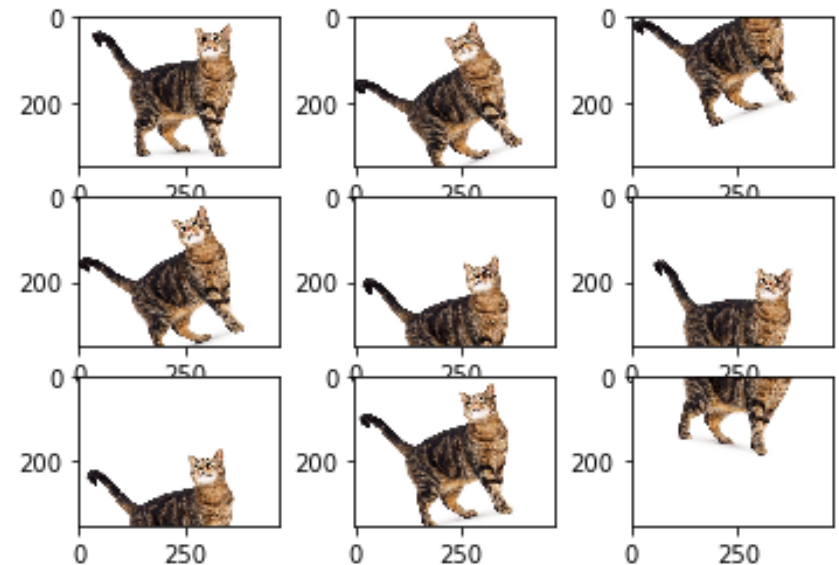
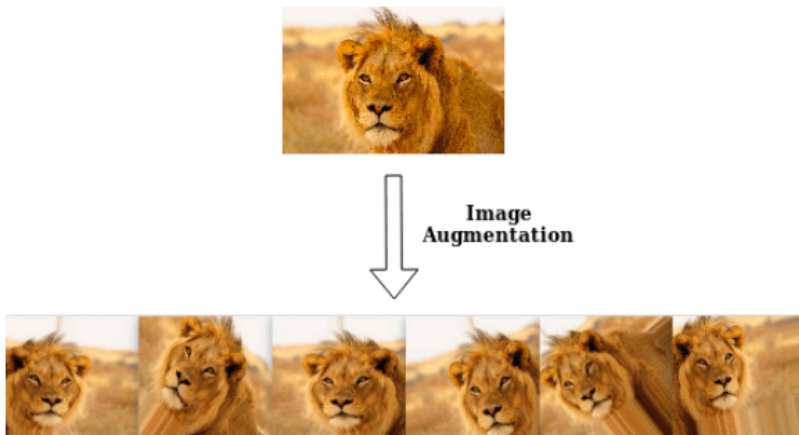
- ▶ p.182-p.192의 Dogs vs. Cats 프로그램을 실행하고 다음 그림을 사용하여 설명하시오.
- ▶ layer 수를 변경하고 매개변수를 조정하여 얻은 결과를 비교하여 설명하시오.



5.2 Training a convnet from scratch on a small dataset

5.2.5 Using data augmentation

- ▶ **Data augmentation** - generating more training data via a number of random transformations
- ▶ expose the model to more aspects of the data and **generalize** better
- ▶ **ImageDataGenerator** instance - number of random transformations




5.2 Training a convnet from scratch on a small dataset

5.2.5 Using data augmentation

Listing 5.11 Setting up a data augmentation configuration via ImageDataGenerator

```
datagen=ImageDataGenerator(rotation_range=40 #degrees  
    width_shift_range=0.2,height_shift_range=0.2,  
    shear_range=0.2, zoom_range=0.2,  
    horizontal_flip=True, fill_mode='nearest')
```



- ▶ These are just a few of the options available (for more, see the Keras documentation):
 - **rotation_range** - degrees (0–180), a range within which to randomly rotate pictures.
 - **width_shift** and **height_shift** - ranges (as a fraction of total width or height) within which to randomly translate pictures vertically or horizontally.
 - **shear_range** - randomly applying shearing transformations.
 - **zoom_range** - randomly zooming inside pictures.
 - **horizontal_flip** - randomly flipping half the images horizontally—relevant when there are no assumptions of horizontal asymmetry (for example, real-world pictures).



5.2 Training a convnet from scratch on a small dataset

5.2.5 Using data augmentation

- **fill_mode** is the strategy used for filling in newly created pixels, which can appear after a rotation or a width/height shift. {"constant", "nearest", "reflect"}.



NEAREST



WRAP



REFLECT



CONSTANT



5.2 Training a convnet from scratch on a small dataset

5.2.5 Using data augmentation

Listing 5.12 Displaying some randomly augmented training images

```
from keras.preprocessing import image # 이미지 전처리 유틸리티 모듈
fnames = sorted([os.path.join(train_cats_dir, fname) for
                    fname in os.listdir(train_cats_dir)])
img_path = fnames[3] # 증식할 이미지 선택
# 이미지를 읽고 크기 변경
img = image.load_img(img_path, target_size=(150, 150))
# (150, 150, 3) 크기의 넘파이 배열로 변환,[:, :, 0:3] 반환
x = image.img_to_array(img)
x = x.reshape((1,)+x.shape) # (1, 150, 150, 3) 크기로 변환
# flow() 메서드는 랜덤하게 변환된 이미지의 배치를 생성
# 무한 반복되기 때문에 어느 지점에서 중지해야 합니다!
i = 0
# flow-이미지를 배치 단위로 가져옴
for batch in datagen.flow(x, batch_size=1):
    plt.figure(i)
    imgplot = plt.imshow(image.array_to_img(batch[0]))
    i += 1
    if i % 4 == 0:
        break
plt.show()
```



cat.100.jpg



Figure 5.11 Generation of cat pictures via random data augmentation

5.2 Training a convnet from scratch on a small dataset

5.2.5 Using data augmentation

- ▶ **data-augmentation** - never produce the same input twice.
- ▶ **overfitting** - remix existing inputs are **still heavily intercorrelated**
- ▶ add a **Dropout** layer to your model, right before the densely connected classifier

Listing 5.13 Defining a new convnet that includes dropout

```
model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu',
                        input_shape=(150, 150, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Flatten())
model.add(layers.Dropout(0.5)) # Flatten 다음, FCN 전
model.add(layers.Dense(512, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))

model.compile(loss='binary_crossentropy',
              optimizer=optimizers.RMSprop(lr=1e-4), metrics=['acc'])
```

5.2 Training a convnet from scratch on a small dataset

Listing 5.14 Training the convnet using data-augmentation

```
train_datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=40,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,)
test_datagen = ImageDataGenerator(rescale=1./255) # 검증 데이터는 증식하지 않음
train_generator = train_datagen.flow_from_directory(
    train_dir, # 타겟 디렉터리
    target_size=(150, 150), # 150 × 150 크기로 바꿉니다
    batch_size=32, # data augmented
    class_mode='binary') # 이진 레이블
validation_generator = test_datagen.flow_from_directory(
    validation_dir,
    target_size=(150, 150),
    batch_size=32,
    class_mode='binary')
history = model.fit_generator(
    train_generator,
    steps_per_epoch=100,
    epochs=100,
    validation_data=validation_generator,
    validation_steps=50)
```


5.2 Training a convnet from scratch on a small dataset

Listing 5.15 Saving the model generators

```
model.save('cats_and_dogs_small_2.h5')
```

- ▶ data augmentation and dropout - no longer overfitting: the training curves are closely tracking the validation curves.
- ▶ accuracy of **82%**, a **15%** **relative** improvement over the non-regularized model.
- ▶ use a pretrained model




Figure 5.12 Training and validation accuracy with data augmentation




Figure 5.13 Training and validation loss with data augmentation

5.3 Using a pretrained convnet

- ▶ **pretrained network** - previously trained on a **large dataset**, typically on a large-scale image-classification task
- ▶ **large and general enough dataset** - generic model, useful for many different computer-vision problems
- ▶ train a network on **ImageNet** (1000 of classes, 1.4 million of images)
- ▶ a key **advantage of deep learning** - portability of learned features across different problems, very effective for **small-data problems**.

5.3 Using a pretrained convnet

- ▶ **ImageNet** contains many animal classes, including different species of **cats and dogs** – perform well on the dogs-versus-cats classification problem.
- ▶ **VGG16 convnet** architecture for ImageNet - by Karen Simonyan and Andrew Zisserman in 2014, easy to understand without introducing any new concepts
- ▶ Previous works - **VGG**, **ResNet**, **Inception**, **Inception-ResNet**, **Xception**
- ▶ There are two ways to use a pretrained network: **feature extraction** and **fine-tuning**.

5.3 Using a pretrained convnet

5.3.1 Feature extraction

- ▶ **Feature extraction** - learned by a previous network to extract interesting features from new samples. These features are then run through a new classifier.
- ▶ convnets used for image classification comprise two parts:
 - *convolutional base* - series of pooling and convolution layers
 - *densely connected classifier*

5.3 Using a pretrained convnet

5.3.1 Feature extraction

▶ Running the new data through it, and training a new classifier on top of the output (figure 5.14).

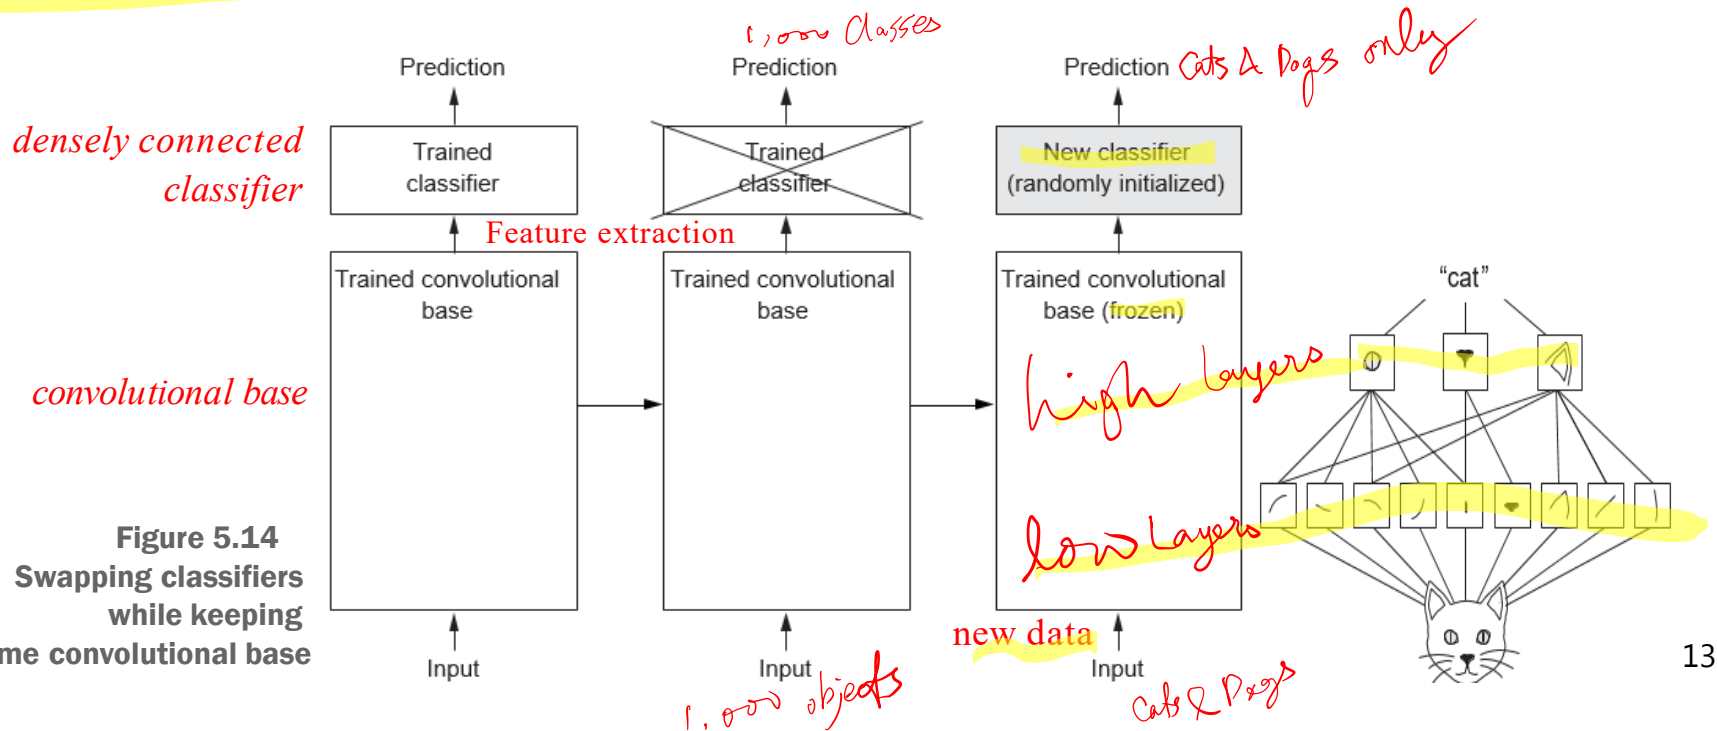


Figure 5.14
Swapping classifiers
while keeping
the same convolutional base

○○○ 5.3 Using a pretrained convnet ○○○

5.3.1 Feature extraction

- ▶ Reuse the densely connected classifier? should be avoided.
- ▶ The feature maps of a convnet are presence maps of **generic concepts** over a picture.
- ▶ level of generality - depth of the layer in the model.
 - low layers - highly **generic** feature maps (such as visual edges, colors, and textures)
 - high layers - more-**abstract** concepts (such as “cat ear” or “dog eye”)

5.3 Using a pretrained convnet

5.3.1 Feature extraction

- ▶ **VGG16** network trained on ImageNet - train a **dogs-versus-cats**
- ▶ Import it from the `keras.applications` module.
- ▶ Here's the list of image-classification models (all pretrained on the ImageNet dataset) that are available as part of `keras.applications`:
 - Xception
 - Inception V3
 - ResNet50
 - **VGG16**
 - VGG19
 - MobileNet

5.3 Using a pretrained convnet

5.3.1 Feature extraction

Listing 5.16 Instantiating the VGG16 convolutional base

```
from keras.applications import VGG16

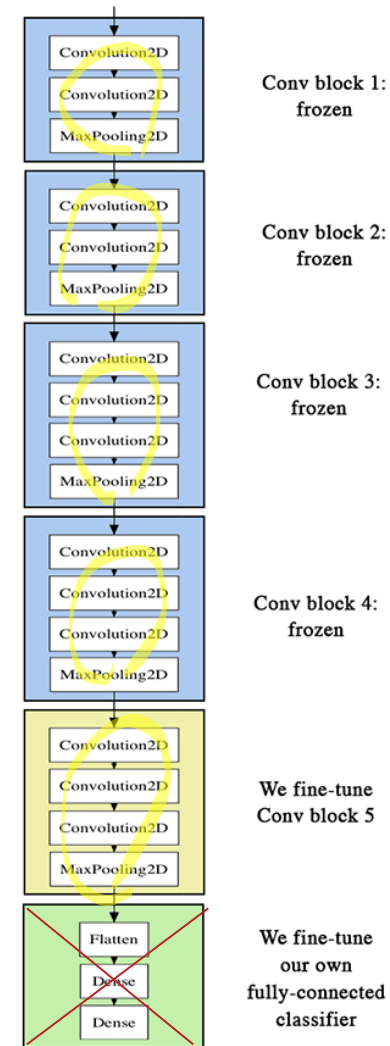
conv_base = VGG16(weights='imagenet',
                   include_top=False, # densely connected classifier
                   input_shape=(150, 150, 3)) # optional
```

- ▶ Three arguments to the constructor:
 - `weights` - weights initialization
 - `include_top` - densely connected classifier on top of the network.
 - By default, 1,000 classes from ImageNet. For the two classes of cat and dog, don't include it.
 - `input_shape` - shape of the image tensors

5.3 Using a pretrained convnet

```
>>> conv_base.summary()
```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(None, 150, 150, 3)	0
block1_conv1 (Convolution2D)	(None, 150, 150, 64)	1792
block1_conv2 (Convolution2D)	(None, 150, 150, 64)	36928
block1_pool (MaxPooling2D)	(None, 75, 75, 64)	0
block2_conv1 (Convolution2D)	(None, 75, 75, 128)	73856
block2_conv2 (Convolution2D)	(None, 75, 75, 128)	147584
block2_pool (MaxPooling2D)	(None, 37, 37, 128)	0
block3_conv1 (Convolution2D)	(None, 37, 37, 256)	295168
block3_conv2 (Convolution2D)	(None, 37, 37, 256)	590080
block3_conv3 (Convolution2D)	(None, 37, 37, 256)	590080
block3_pool (MaxPooling2D)	(None, 18, 18, 256)	0
block4_conv1 (Convolution2D)	(None, 18, 18, 512)	1180160
block4_conv2 (Convolution2D)	(None, 18, 18, 512)	2359808
block4_conv3 (Convolution2D)	(None, 18, 18, 512)	2359808
block4_pool (MaxPooling2D)	(None, 9, 9, 512)	0
block5_conv1 (Convolution2D)	(None, 9, 9, 512)	2359808
block5_conv2 (Convolution2D)	(None, 9, 9, 512)	2359808
block5_conv3 (Convolution2D)	(None, 9, 9, 512)	2359808
block5_pool (MaxPooling2D)	(None, 4, 4, 512)	0



5.3 Using a pretrained convnet

5.3.1 Feature extraction

- ▶ The final feature map has shape $(4, 4, 512)$. That's the **feature** on top of densely connected classifier.
- ▶ two ways to proceed:
 - Running the convolutional base over your dataset → recording its output to a Numpy array (features) → input to densely connected classifier - running the convolutional base once for every input image **without data augmentation**.
 - Running the whole thing end to end on the input data **with data augmentation**