

## 5.3 Using a pretrained convnet

FAST FEATURE EXTRACTION WITHOUT DATA AUGMENTATION

**Listing 5.17** Extracting features using the pretrained convolutional base

```
import os
import numpy as np
from keras.preprocessing.image import ImageDataGenerator
base_dir = './datasets/cats_and_dogs_small'
train_dir = os.path.join(base_dir, 'train')
validation_dir = os.path.join(base_dir, 'validation')
test_dir = os.path.join(base_dir, 'test')
datagen = ImageDataGenerator(rescale=1./255) # WITHOUT DATA AUGMENTATION
def extract_features(directory, sample_count):
    features = np.zeros(shape=(sample_count, 4, 4, 512))
    labels = np.zeros(shape=(sample_count))
    generator = datagen.flow_from_directory(
        directory,          target_size=(150, 150),
        batch_size=20,      class_mode='binary')
    i = 0
    for inputs_batch, labels_batch in generator:
        features_batch = conv_base.predict(inputs_batch) # conv_base = vgg16
        features[i * 20 : (i + 1) * 20] = features_batch
        labels[i * 20 : (i + 1) * 20] = labels_batch
        i += 1
        if i * 20 >= sample_count: # WITHOUT DATA AUGMENTATION
            break
    return features, labels
train_features, train_labels = extract_features(train_dir, 2000)
validation_features, validation_labels = extract_features(validation_dir, 1000)
test_features, test_labels = extract_features(test_dir, 1000)
```

## 5.3 Using a pretrained convnet

► The extracted **features** are currently of shape (samples, 4, 4, 512) → densely connected classifier (samples, 8192):

```
train_features = np.reshape(train_features, (2000, 4 * 4 * 512))
validation_features = np.reshape(validation_features, (1000, 4 * 4 * 512))
test_features = np.reshape(test_features, (1000, 4 * 4 * 512))
```

*flatten*

### Listing 5.18 Defining and training the densely connected classifier

```
from keras import models
from keras import layers
from keras import optimizers
model = models.Sequential()
model.add(layers.Dense(256, activation='relu', input_dim=4 * 4 * 512))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(1, activation='sigmoid'))
model.compile(optimizer=optimizers.RMSprop(lr=2e-5),
              loss='binary_crossentropy',
              metrics=['acc'])
history = model.fit(train_features, train_labels,
                    epochs=30,
                    batch_size=20,
                    validation_data=(validation_features, validation_labels))
```

► Training is very fast, because you only have to deal with two Dense layers—an epoch takes less than one second even on CPU.

## 5.3 Using a pretrained convnet

### Listing 5.19 Plotting the results

```
import matplotlib.pyplot as plt
acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(len(acc))
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```




Figure 5.15 Training and validation accuracy for simple feature extraction




Figure 5.16 Training and validation loss for simple feature extraction

► You reach a validation accuracy of about 90%. But the plots also indicate that you're **overfitting** almost from the start—despite using **dropout** with a fairly large rate. That's because this technique **doesn't use data augmentation**, which is essential for preventing overfitting with **small image datasets**.

## 5.3 Using a pretrained convnet

### FEATURE EXTRACTION WITH DATA AUGMENTATION

- ▶ data augmentation during training - much slower and more expensive
- ▶ extending the `conv_base` model and running it end to end on the inputs

▶**NOTE** This technique is so expensive that you should only attempt it if you have access to a GPU—it's absolutely intractable on CPU. If you can't run your code on GPU, then the previous technique is the way to go

### Listing 5.20 Adding a densely connected classifier on top of the convolutional base

```
from keras import models
from keras import layers
model = models.Sequential()
model.add(conv_base)
model.add(layers.Flatten())
model.add(layers.Dense(256, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

```
>>> model.summary()
```

Layer (type)	Output Shape	Param #
vgg16 (Model)	(None, 4, 4, 512)	14714688
flatten_1 (Flatten)	(None, 8192)	0
dense_1 (Dense)	(None, 256)	2097408
dense_2 (Dense)	(None, 1)	257

```
=====  
Total params: 16,812,353  
Trainable params: 16,812,353  
Non-trainable params: 0
```

## 5.3 Using a pretrained convnet

### FEATURE EXTRACTION WITH DATA AUGMENTATION

- ▶ **Freezing** - freeze the convolutional base, preventing weights from being updated during training
- ▶ If you don't do this, then the representations that were previously learned by the convolutional base will be modified during training.

- ▶ In Keras, you freeze a network by setting its **trainable** attribute to **False**:

```
>>> print('This is the number of trainable weights'
        'before freezing the conv base:', len(model.trainable_weights))
This is the number of trainable weights before freezing the conv base: 30
>>> conv_base.trainable = False
>>> print('This is the number of trainable weights '
        'after freezing the conv base:', len(model.trainable_weights))
This is the number of trainable weights after freezing the conv base: 4
```



- ▶ **2 Dense layers** that you added will be trained - That's a total of **4 weight tensors**: two per layer (the **main weight matrix** and **the bias vector**).  $(13+2)*2=30$  weight tensors for whole model.
- ▶ Now you can start training your model, with the **same data-augmentation** configuration that you used in the previous example.

## 5.3 Using a pretrained convnet

### Listing 5.21 Training the model end to end with a frozen convolutional base

```
from keras.preprocessing.image import ImageDataGenerator
train_datagen = ImageDataGenerator(
    rescale=1./255,      rotation_range=20,
    width_shift_range=0.1, height_shift_range=0.1,
    shear_range=0.1,     zoom_range=0.1,
    horizontal_flip=True, fill_mode='nearest')
test_datagen=ImageDataGenerator(rescale=1./255) # No augmented
train_generator = train_datagen.flow_from_directory(
    train_dir, # 타겟 디렉터리
    target_size=(150, 150), # 150 × 150로 변경
    batch_size=20,
    class_mode='binary') # 이진 레이블
validation_generator = test_datagen.flow_from_directory(
    validation_dir,
    target_size=(150, 150),
    batch_size=20,      class_mode='binary')
model.compile(loss='binary_crossentropy',
    optimizer=optimizers.RMSprop(lr=2e-5),
    metrics=['acc']) # augmented!
history = model.fit_generator(
    train_generator, # 2000/100=20 data augmentations
    steps_per_epoch=100,      epochs=30,
    validation_data=validation_generator,
    validation_steps=50,
    verbose=2) # 진행 막대(progress bar)가 나오지 않도록 설정
```

▶ validation accuracy of 96% - better the small convnet trained from scratch



## 5.3 Using a pretrained convnet

### 5.3.2 Fine-tuning

- ▶ ***fine-tuning*** - slightly adjusts the **more abstract** representations of the model being reused, in order to make them more relevant for the problem at hand.
- ▶ fine-tune the **top layers** of the convolutional base
- ▶ The steps for fine-tuning a network are as follow:
  - 1 **Add** your custom network on top of an already-trained base network.
  - 2 **Freeze** the base network.
  - 3 **Train** the part you added.
  - 4 **Unfreeze** some layers in the base network.
  - 5 Jointly **train both** these layers and the part you added.

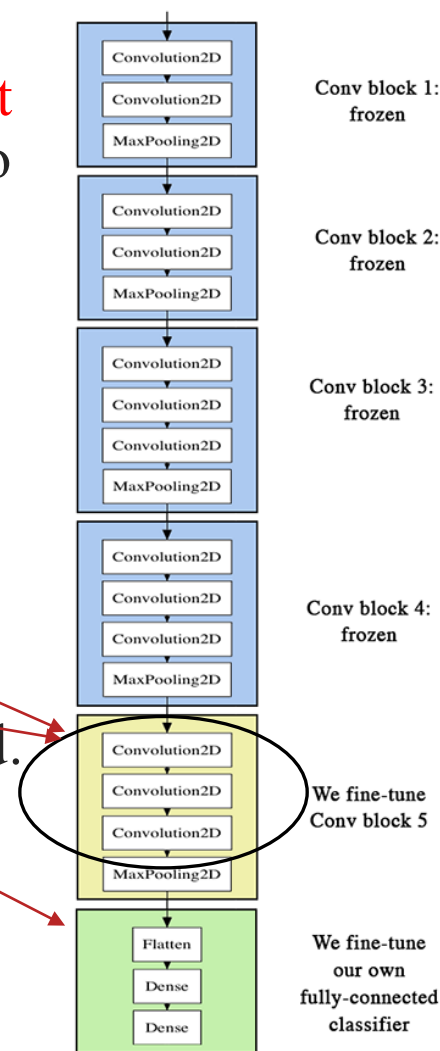


Figure 5.19 Fine-tuning the last convolutional block of the VGG16 network

## 5.3 Using a pretrained convnet

- ▶ **fine-tune** - the layers `block5_conv1`, `block5_conv2`, and `block5_conv3` should be **trainable**.
- ▶ Why not fine-tune more layers? Why not fine-tune the entire convolutional base? consider the following:

- Earlier layers - **more-generic**, reusable features
- Higher layers - **more-specialized** features.
- The more parameters you're training, the more you're at risk of **overfitting**.
- The convolutional base has 15 million parameters, so it would be **risky** to attempt to train it on your **small dataset**.

### Listing 5.22 Freezing all layers up to a specific one

```
conv_base.trainable = True # T,T,...,T
set_trainable = False
for layer in conv_base.layers:
    if layer.name == 'block5_conv1':
        set_trainable = True
    if set_trainable: # F,F,...,F,T,T,T
        layer.trainable = True
    else: # freeze before block5_conv1
        layer.trainable = False
```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(None, 150, 150, 3)	0
block1_conv1 (Conv2D)	(None, 150, 150, 64)	1792
block1_conv2 (Conv2D)	(None, 150, 150, 64)	36928
block1_pool (MaxPooling2D)	(None, 75, 75, 64)	0
block2_conv1 (Conv2D)	(None, 75, 75, 128)	73856
block2_conv2 (Conv2D)	(None, 75, 75, 128)	147584
block2_pool (MaxPooling2D)	(None, 37, 37, 128)	0
block3_conv1 (Conv2D)	(None, 37, 37, 256)	295168
block3_conv2 (Conv2D)	(None, 37, 37, 256)	590080
block3_conv3 (Conv2D)	(None, 37, 37, 256)	590080
block3_pool (MaxPooling2D)	(None, 18, 18, 256)	0
block4_conv1 (Conv2D)	(None, 18, 18, 512)	1180160
block4_conv2 (Conv2D)	(None, 18, 18, 512)	2359808
block4_conv3 (Conv2D)	(None, 18, 18, 512)	2359808
block4_pool (MaxPooling2D)	(None, 9, 9, 512)	0
block5_conv1 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv2 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv3 (Conv2D)	(None, 9, 9, 512)	2359808
block5_pool (MaxPooling2D)	(None, 4, 4, 512)	0
Total params: 14,714,688		
Trainable params: 0		
Non-trainable params: 14,714,688		



## 5.3 Using a pretrained convnet

### Listing 5.23 Fine-tuning the model

```
model.compile(loss='binary_crossentropy',
              optimizer=optimizers.RMSprop(lr=1e-5),
              metrics=['acc'])
history = model.fit_generator(
    train_generator,
    steps_per_epoch=100,
    epochs=100,
    validation_data=validation_generator,
    validation_steps=50)
```

## 5.3 Using a pretrained convnet

### Listing 5.24 Smoothing the plots

```
def smooth_curve(points, factor=0.8):
    smoothed_points = []
    for point in points:
        if smoothed_points:
            previous = smoothed_points[-1]
            smoothed_points.append(previous * factor + point * (1 - factor))
        else:
            smoothed_points.append(point)
    return smoothed_points

plt.plot(epochs,
         smooth_curve(acc), 'bo', label='Smoothed training acc')
plt.plot(epochs,
         smooth_curve(val_acc), 'b', label='Smoothed validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs,
         smooth_curve(loss), 'bo', label='Smoothed training loss')
plt.plot(epochs,
         smooth_curve(val_loss), 'b', label='Smoothed validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```

- ▶ The validation accuracy curve from about 96% to above 97%
- ▶ Note that the loss curve doesn't show any real improvement (in fact, it's deteriorating)




Figure 5.22 Smoothed curves for training and validation accuracy for fine-tuning




Figure 5.23 Smoothed curves for training and validation loss for fine-tuning

## 5.3 Using a pretrained convnet

- ▶ You can now finally evaluate this model on the **test data**:

```
test_generator  
=test_datagen.flow_from_directory(  
    test_dir,  
    target_size=(150, 150),  
    batch_size=20,  
    class_mode='binary')
```

- ▶ test accuracy of **92%** - In the original **Kaggle competition** around this dataset, this would have been one of the **top results** using only a small fraction of the training data available (about **10%**). There is a huge difference between being able to train on **20,000** samples compared to **2,000** samples!

## 5.3 Using a pretrained convnet

### ▶ 실습

다음 convolution block들을 unfreeze하고 fine tuning 하는 프로그램과 그 결과를 ppt 10 page와 비교하시오.

1) `block5_conv3`

2) `block5_conv3, block5_conv2`