

## 5.3 Using a pretrained convnet

► **fine-tune** - the layers `block5_conv1`, `block5_conv2`, and `block5_conv3` should be **trainable**.

► Why not fine-tune more layers? Why not fine-tune the entire convolutional base? consider the following:

- Earlier layers - **more-generic**, reusable features
- Higher layers - **more-specialized** features.
- The more parameters you're training, the more you're at risk of **overfitting**.
- The convolutional base has 15 million parameters, so it would be **risky** to attempt to train it on your **small dataset**.

### Listing 5.22 Freezing all layers up to a specific one

```
conv_base.trainable = True # T,T,...,T
set_trainable = False
for layer in conv_base.layers:
    if layer.name == 'block5_conv1':
        set_trainable = True
    if set_trainable: # F,F,...,F,T,T,T
        layer.trainable = True
    else: # freeze before block5_conv1
        layer.trainable = False
```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(None, 150, 150, 3)	0
block1_conv1 (Conv2D)	(None, 150, 150, 64)	1792
block1_conv2 (Conv2D)	(None, 150, 150, 64)	36928
block1_pool (MaxPooling2D)	(None, 75, 75, 64)	0
block2_conv1 (Conv2D)	(None, 75, 75, 128)	73856
block2_conv2 (Conv2D)	(None, 75, 75, 128)	147584
block2_pool (MaxPooling2D)	(None, 37, 37, 128)	0
block3_conv1 (Conv2D)	(None, 37, 37, 256)	295168
block3_conv2 (Conv2D)	(None, 37, 37, 256)	590080
block3_conv3 (Conv2D)	(None, 37, 37, 256)	590080
block3_pool (MaxPooling2D)	(None, 18, 18, 256)	0
block4_conv1 (Conv2D)	(None, 18, 18, 512)	1180160
block4_conv2 (Conv2D)	(None, 18, 18, 512)	2359808
block4_conv3 (Conv2D)	(None, 18, 18, 512)	2359808
block4_pool (MaxPooling2D)	(None, 9, 9, 512)	0
block5_conv1 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv2 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv3 (Conv2D)	(None, 9, 9, 512)	2359808
block5_pool (MaxPooling2D)	(None, 4, 4, 512)	0
Total params: 14,714,688		
Trainable params: 0		
Non-trainable params: 14,714,688		

## 5.3 Using a pretrained convnet

### Listing 5.23 Fine-tuning the model

```
model.compile(loss='binary_crossentropy',
              optimizer=optimizers.RMSprop(lr=1e-5),
              metrics=['acc'])
history = model.fit_generator(
    train_generator,
    steps_per_epoch=100,
    epochs=100,
    validation_data=validation_generator,
    validation_steps=50)
```

## 5.3 Using a pretrained convnet

### Listing 5.24 Smoothing the plots

```
def smooth_curve(points, factor=0.8):
    smoothed_points = []
    for point in points:
        if smoothed_points:
            previous = smoothed_points[-1]
            smoothed_points.append(previous * factor + point * (1 - factor))
        else:
            smoothed_points.append(point)
    return smoothed_points

plt.plot(epochs,
         smooth_curve(acc), 'bo', label='Smoothed training acc')
plt.plot(epochs,
         smooth_curve(val_acc), 'b', label='Smoothed validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs,
         smooth_curve(loss), 'bo', label='Smoothed training loss')
plt.plot(epochs,
         smooth_curve(val_loss), 'b', label='Smoothed validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```

- ▶ The validation accuracy curve from about 96% to above 97%
- ▶ Note that the loss curve doesn't show any real improvement (in fact, it's deteriorating)




Figure 5.22 Smoothed curves for training and validation accuracy for fine-tuning




Figure 5.23 Smoothed curves for training and validation loss for fine-tuning

## 5.3 Using a pretrained convnet

- ▶ You can now finally evaluate this model on the **test data**:

```
test_generator  
=test_datagen.flow_from_directory(  
    test_dir,  
    target_size=(150, 150),  
    batch_size=20,  
    class_mode='binary')
```

- ▶ test accuracy of **92%** - In the original **Kaggle competition** around this dataset, this would have been one of the **top results** using only a small fraction of the training data available (about **10%**). There is a huge difference between being able to train on **20,000** samples compared to **2,000** samples!

## 5.3 Using a pretrained convnet

### ▶ 실습

다음 convolution block들을 unfreeze하고 fine tuning 하는 프로그램과 그 결과를 ppt 58 page와 비교하시오.

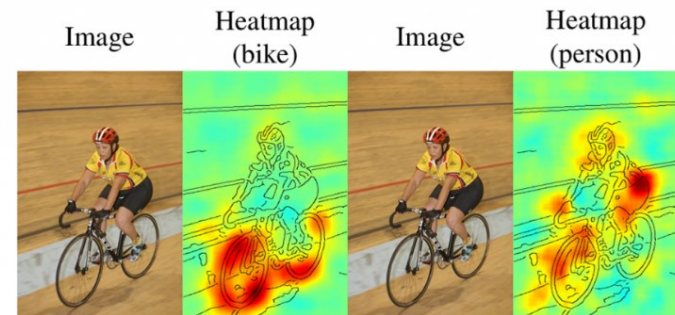

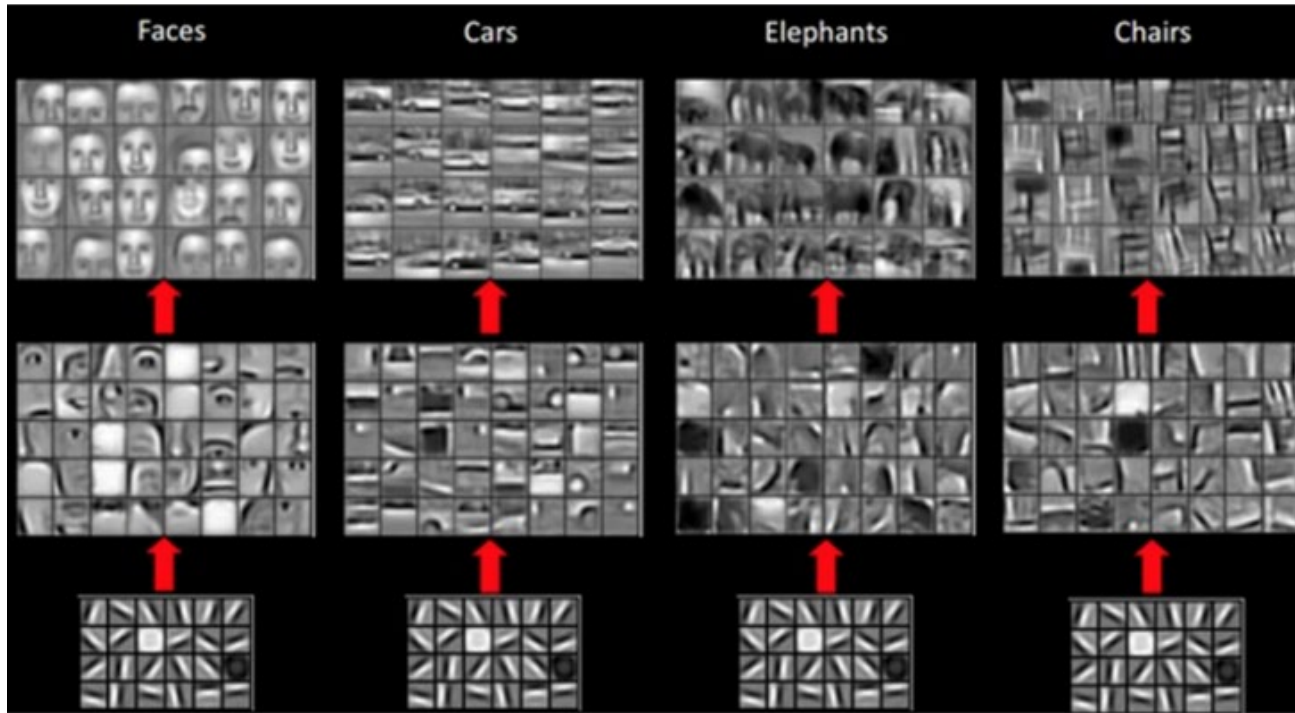
1) `block5_conv3`

2) `block5_conv3, block5_conv2`

## 5.4 Visualizing what convnets learn

- ▶ “**black boxes**”: learning representations that are difficult to extract and present in a **human-readable** form
- ▶ **convnets** are highly amenable to visualization - *representations of **visual concepts***
- ▶ Since 2013, a wide array of techniques have been developed for **visualizing** and **interpreting** these representations:
  - Visualizing **intermediate convnet** outputs (**intermediate activations**) — Understanding **convnet layers** transform their input, and **individual convnet filters**
  - Visualizing **convnets filters** — understanding precisely what **visual pattern** or **concept** each filter in a convnet is receptive to.
  - Visualizing **heatmaps** of class activation in an image—identified as belonging to a given class, localize objects in images.

# 5.4 Visualizing what convnets learn



Explaining classification "bike"

Explaining classification "person"

# 5.4 Visualizing what convnets learn

## 5.4.1 Visualizing intermediate activations

- ▶ Visualizing intermediate activations - **convolution** and **pooling** layers
- ▶ How an **input** is decomposed into the different filters learned by the network.
- ▶ visualize **feature maps** - plotting the contents of every channel as a **2D image**
- ▶ Let's start by loading the model that you saved in section 5.2:

```
>>> from keras.models import load_model
>>> model= load_model('cats_and_dogs_small_2.h5')
>>> model.summary() # As a reminder
```

Layer (type)	Output Shape	Param #
conv2d_5 (Conv2D)	(None, 148, 148, 32)	896
max_pooling2d_5 (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_6 (Conv2D)	(None, 72, 72, 64)	18496
max_pooling2d_6 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_7 (Conv2D)	(None, 34, 34, 128)	73856
max_pooling2d_7 (MaxPooling2D)	(None, 17, 17, 128)	0
conv2d_8 (Conv2D)	(None, 15, 15, 128)	147584
max_pooling2d_8 (MaxPooling2D)	(None, 7, 7, 128)	0
flatten_2 (Flatten)	(None, 6272)	0
dropout_1 (Dropout)	(None, 6272)	0
dense_3 (Dense)	(None, 512)	3211776
dense_4 (Dense)	(None, 1)	513




## 5.4 Visualizing what convnets learn

### Listing 5.25 Preprocessing a single image

```
img_path =  
    './datasets/cats_and_dogs_small/test/cats/cat.1700.jpg'  
  
# 이미지를 4D 텐서로 변경합니다  
from keras.preprocessing import image  
import numpy as np  
img = image.load_img(img_path, target_size=(150, 150))  
img_tensor = image.img_to_array(img) # (150, 150, 3) # 컬러 이미지  
img_tensor = np.expand_dims(img_tensor, axis=0) # 맨 앞차원에 추가  
# img_tensor = img_tensor.reshape((1,) + img_tensor.shape)  
# 모델이 훈련될 때 입력에 적용한 전처리 방식  
img_tensor /= 255.  
  
print(img_tensor.shape) # (1, 150, 150, 3)
```

### Listing 5.25 Preprocessing a single image

```
import matplotlib.pyplot as plt  
  
plt.imshow(img_tensor[0])  
plt.show()
```



# 5.4 Visualizing what convnets learn

- ▶ outputs the activations of all **convolution** and **pooling** layers.
- ▶ **Keras class Model** – an **input** tensor (or list of input tensors) and an **output** tensor (or list of output tensors)

## **Listing 5.27 Instantiating a model from an input tensor and a list of output tensors**

```
from keras import models
```

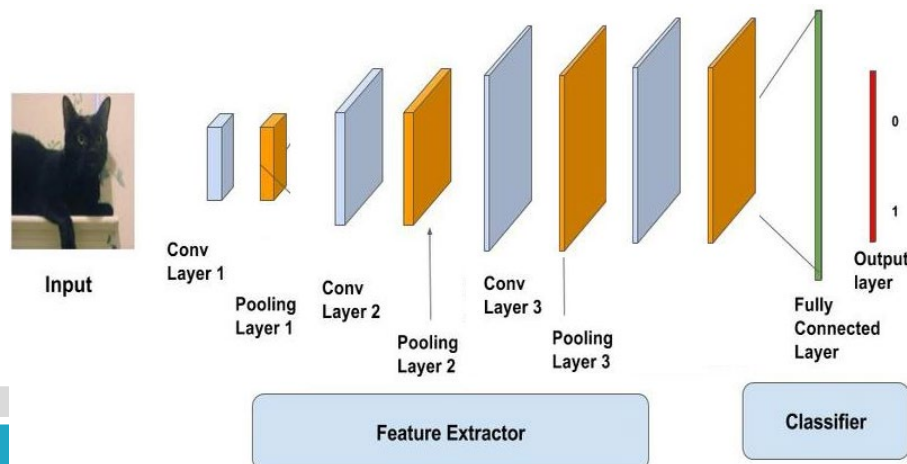
```
# Extracts the outputs of the top eight layers
```

```
layer_outputs = [layer.output for layer in  
                  model.layers[:8]]  
# model= load_model('cats_and_dogs_small_2.h5')
```

```
# Creates a model that will return these outputs, given the model input
```

```
activation_model = models.Model(inputs=model.input, #cat image  
                                outputs=layer_outputs) # 8 output layers
```

- ▶ **multi-output** model - **one input** and **eight outputs**, one output per layer activation



## 5.4 Visualizing what convnets learn

### Listing 5.28 Running the model in predict mode

# Returns a list of 8 Numpy arrays: one array per layer activation

```
activations = activation_model.predict(img_tensor) #cat image
```

- ▶ activation of the first convolution layer for the cat image input:

```
>>> first_layer_activation = activations[0]
>>> print(first_layer_activation.shape)
(1, 148, 148, 32)
```

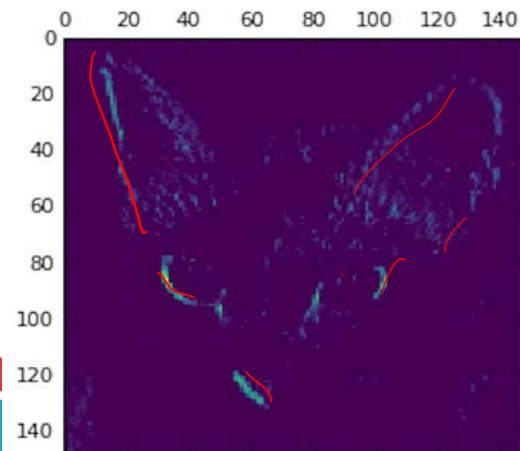
- ▶ It's a  $148 \times 148$  feature map with 32 channels. Let's try plotting the **fourth channel** of the activation of the **first layer** of the original model (see figure 5.25).

### Listing 5.29 Visualizing the fourth channel

```
import matplotlib.pyplot as plt
plt.matshow(first_layer_activation[0, :, :, 4], cmap='viridis')
```

- ▶ This channel appears to encode a **diagonal edge** detector.

Figure 5.25 Fourth channel of the activation of the first layer on the test cat picture



## 5.4 Visualizing what convnets learn

- ▶ Let's try **the seventh** channel — the specific filters learned by convolution layers aren't deterministic.

### Listing 5.30 Visualizing the seventh channel

```
plt.matshow(first_layer_activation[0, :, :, 7],  
            cmap='viridis')
```

- ▶ This one looks like a “**bright green dot**” detector, useful to encode cat eyes.
- ▶ extract and plot every channel in each of the eight activation maps




Figure 5.26 Seventh channel of the activation of the first layer on the test cat picture