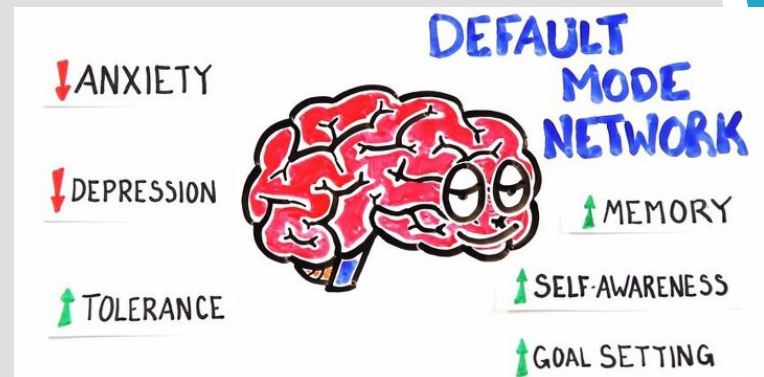



6장 *Deep learning for text and sequences*

“default mode network”





This chapter covers



- Preprocessing text data into useful representations – **tokens to vectors**
- Working with **recurrent neural networks**
- Using 1D convnets for **sequence processing**

○○○ *This chapter covers* ○○○

- ▶ sequences of **word**, **timeseries**, and sequence **data** in general
 - ▶ *recurrent neural networks* and *1D convnets*
 - ▶ Applications of these algorithms include the following:
 - Document classification - identifying the **topic of an article** or the **author of a book**
 - Sequence-to-sequence learning - **English sentence into French**
 - Sentiment analysis - classifying the sentiment of **tweets** or **movie reviews** as positive or negative
 - Timeseries forecasting - predicting the **future weather** given recent weather data
1. **sentiment** analysis on the IMDB dataset
 2. **temperature** forecasting.

6.1 Working with text data

- ▶ **natural-language** understanding - document classification(topic), sentiment analysis, author identification, and even question-answering (QA)
- ▶ Deep learning for natural-language processing is **pattern recognition** applied to words, sentences, and paragraphs
- ▶ **Vectorizing text** is the process of transforming **text** into **numeric tensors**.
 - ▶ Segment text into **words**, and transform each word into a **vector**.
 - ▶ Segment text into **characters**, and transform each character into a **vector**.
 - ▶ Extract **n-grams** of words or characters, and transform each n-gram into a **vector**.
- ▶ **N-grams** are overlapping groups of **multiple consecutive words or characters**.
- ▶ **tokens** - break down text (**words, characters, or n-grams**)
- ▶ **tokenization** - breaking text into such tokens
- ▶ two major ones: **one-hot encoding** of tokens, and **token embedding**

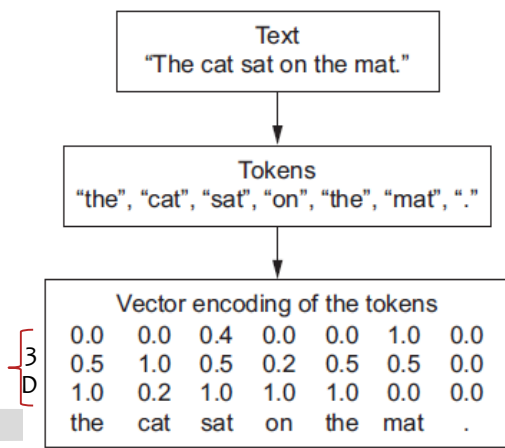


Figure 6.1 From text to tokens to vectors

6.1 Working with text data

Understanding **n-grams** and **bag-of-words (BoW)**

▶ **Word n-grams** are groups of N (or fewer) consecutive words that you can extract from a sentence.

▶ “The cat sat on the mat.” – **set of 2-grams**:

`{"The", "The cat", "cat", "cat sat", "sat",
"sat on", "on", "on the", "the", "the mat", "mat"}` **bag-of-2-grams**

▶ It may also be decomposed into the following set of **3-grams**:

`{"The", "The cat", "cat", "cat sat", "The cat sat",
"sat", "sat on", "on", "cat sat on", "on the", "the", "sat
on the", "the mat", "mat", "on the mat"}` **bag-of-3-grams**

- ▶ Because bag-of-words **isn't an order-preserving tokenization method** (the tokens generated are understood as a **set**, not a sequence, and the general structure of the sentences is lost)
- ▶ unavoidable **feature-engineering** tool when using lightweight, **shallow text-processing models** such as **logistic regression** and **random forests**.

6.1 Working with text data

6.1.1 One-hot encoding of words and characters

- ▶ **One-hot encoding** - turn a token into a vector
- ▶ **IMDB** and **Reuters** examples - done with words
- ▶ a unique **integer index** with every word - **binary vector** of size N (the size of the vocabulary)
- ▶ one-hot encoding can be done at the **character** level

6.1 Working with text data

Listing 6.1 Word-level one-hot encoding (toy example)

```
import numpy as np
samples = ['The cat sat on the mat.', 'The dog ate my homework.']
# 데이터에 있는 모든 토큰의 인덱스를 구축합니다
token_index = {} # dictionary - {key:value}; {'The':1 }
for sample in samples:
    # split() 메서드를 사용해 샘플을 토큰으로 나눕니다.
    for word in sample.split(): # key - word
        if word not in token_index :
            token_index[word] = len(token_index) + 1
            # 인덱스 0은 사용하지 않습니다.
# {'The': 1, 'cat': 2, 'sat': 3, 'on': 4, 'the': 5, 'mat.': 6, 'dog': 7, 'ate': 8, 'my': 9, 'homework.': 10}
# 샘플을 벡터로 변환
max_length = 10
results = np.zeros((len(samples), max_length, max(token_index.values())+1)) # (2,10,11)
for i, sample in enumerate(samples):
    for j, word in list(enumerate(sample.split()))[:max_length]:
        index = token_index.get(word)
        results[i, j, index] = 1.
```

[[0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0.]	The	[[0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0.]	The
[0. 0. 1. 0. 0. 0. 0. 0. 0. 0. 0.]	cat	[0. 0. 0. 0. 0. 0. 0. 1. 0. 0. 0.]	dog
[0. 0. 0. 1. 0. 0. 0. 0. 0. 0. 0.]	sat	[0. 0. 0. 0. 0. 0. 0. 0. 1. 0. 0.]	ate
[0. 0. 0. 0. 1. 0. 0. 0. 0. 0. 0.]	on	[0. 0. 0. 0. 0. 0. 0. 0. 0. 1. 0.]	my
[0. 0. 0. 0. 0. 1. 0. 0. 0. 0. 0.]	the	[0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 1.]	homework.
[0. 0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]	mat.	[0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]	
[0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]		[0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]	
[0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]		[0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]	
[0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]		[0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]	
[0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]		[0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]	

6.1 Working with text data

Listing 6.2 Character-level one-hot encoding (toy example)

```
import string
samples = ['The cat sat on the mat.', 'The dog ate my homework.']
characters = string.printable # 출력 가능한 모든 ASCII 문자, 100개
token_index = dict(zip(characters, range(1, len(characters) + 1)))
max_length = 50

results = np.zeros((len(samples), max_length,
                    max(token_index.values())+1))
for i, sample in enumerate(samples):
    for j, character in enumerate(sample[:max_length]):
        index = token_index.get(character)
        results[i, j, index] = 1.
```

token_index = {'0': 1, '1': 2, '2': 3, '3': 4, '4': 5, '5': 6, '6': 7, '7': 8, '8': 9, '9': 10, 'a': 11, 'b': 12, 'c': 13, 'd': 14, 'e': 15, 'f': 16, 'g': 17, 'h': 18, 'i': 19, 'j': 20, 'k': 21, ..., 'A': 37, 'B': 38, ..., '\x0c': 100}

[[[0. 0. 0. ... 0. 0. 0.] [0. 0. 0. ... 0. 0. 0.] [0. 0. 0. ... 1. ... 0. 0. 0.] ... [0. 0. 0. ... 0. 0. 0.] [0. 0. 0. ... 0. 0. 0.] [0. 0. 0. ... 0. 0. 0.]] 'The cat sat on the mat.'

[[[0. 0. 0. ... 0. 0. 0.] [0. 0. 0. ... 1. ... 0. 0.] [0. 0. 0. ... 0. 0. 0.] ... [0. 0. 0. ... 0. 0. 0.] [0. 0. 0. ... 0. 0. 0.] [0. 0. 0. ... 0. 0. 0.]] 'The dog ate my homework.'

6.1 Working with text data

Listing 6.3 Using Keras for word-level one-hot encoding

```
from keras.preprocessing.text import Tokenizer
samples = ['The cat sat on the mat.', 'The dog ate my homework.']
# 가장 빈도가 높은 1,000개의 단어만 선택하도록 Tokenizer 객체를 만듭니다.
tokenizer = Tokenizer(num_words=1000)
# Turns strings into lists of integer indices by word_index
tokenizer.fit_on_texts(samples) # 입력에 맞게 내부의 word_index를 중복 없이 만드는 함수
# tokenizer.word_index = {'the': 1, 'cat': 2, 'sat': 3, 'on': 4, 'mat': 5, 'dog': 6, 'ate': 7, 'my': 8, 'homework': 9}

# Turns strings into lists of integer indices
sequences = tokenizer.texts_to_sequences(samples)
# Sequences = [[1, 2, 3, 4, 1, 5], [1, 6, 7, 8, 9]]

# directly get the one-hot binary representations.
# Vectorization modes other than one-hot encoding are supported by this tokenizer!
one_hot_results = tokenizer.texts_to_matrix(samples, mode='binary')
# one_hot_results = [[0. 1. 1. 1. 1. 0. 0. 0. 0. 0. ... 0. 0. 0.]
#                    [0. 1. 0. 0. 0. 0. 1. 1. 1. 1. 0. ... 0. 0. 0.]]

word_index = tokenizer.word_index
print('Found %s unique tokens.' % len(word_index))
# Found 9 unique tokens.
```

6.1 Working with text data

- ▶ *one-hot hashing* - vocabulary is too large to handle explicitly
- ▶ hash words into vectors of fixed size with a very lightweight **hashing function**
- ▶ saves memory and allows online encoding of the data
- ▶ *hash collisions*: two different words may end up with the same hash

Listing 6.4 Using Keras for word-level one-hot encoding

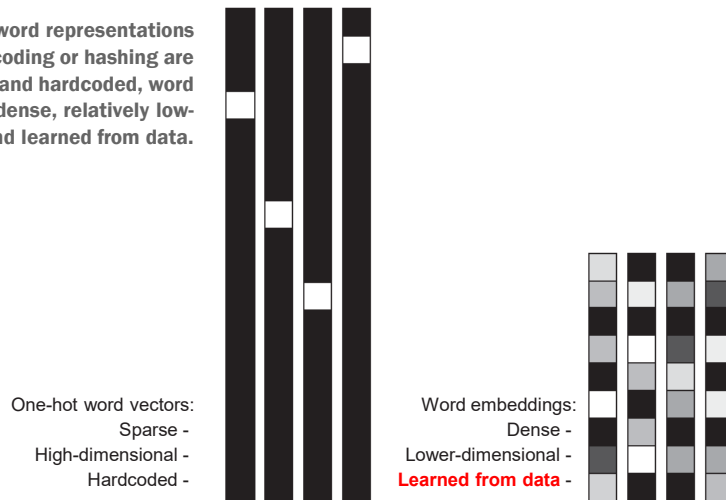
```
samples = ['The cat sat on the mat.', 'The dog ate my homework.']
# 1,000개 이상의 단어가 있다면 hash collisions
dimensionality = 1000
max_length = 10
results = np.zeros((len(samples), max_length, dimensionality))
for i, sample in enumerate(samples):
    for j, word in list(enumerate(sample.split()))[:max_length]:
        # Hashes the word into a random integer index between 0 and 1,000
        index = abs(hash(word)) % dimensionality
        results[i, j, index] = 1.
# in case of dimensionality = 20 →
results[0] =
[[0.0.0.0.0.1.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.] The
[0.0.0.0.1.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.] cat
[0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.1.0.0.0.0.0.0.0.] sat
[0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.1.0.0.0.0.0.0.] on
[0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.1.0.] the
[0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.1.] mat
[0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.]
[0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.]
[0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.]
[0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.0.]
```

6.1 Working with text data

6.1.2 Using word embeddings

- ▶ **one-hot** encoding are binary, sparse, very high-dimensional (20,000-dimensional or greater)
- ▶ **dense word vectors** also called **word embeddings** - **low-dimensional** floating-point vectors in 256-, 512-, or 1,024-dimensional when dealing with very large vocabularies
- ▶ pack more information into far fewer dimensions

Figure 6.2 Whereas word representations obtained from one-hot encoding or hashing are sparse, high-dimensional, and hardcoded, word embeddings are dense, relatively low-dimensional, and learned from data.



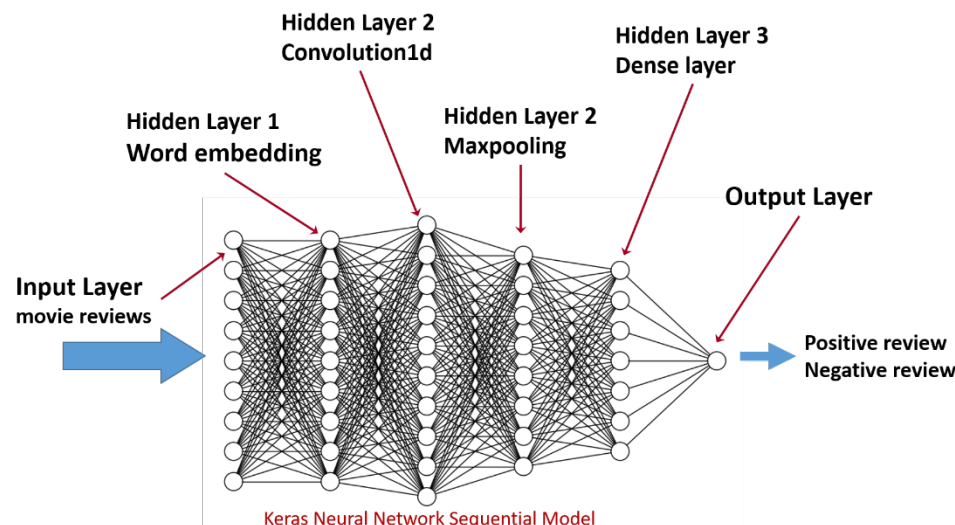

6.1 Working with text data

- ▶ There are two ways to obtain word embeddings:
 - Learn word embeddings jointly with the **main task** you care about (such as **document classification** or **sentiment prediction** → **weights** in a neural network).
 - *pretrained word embeddings* - Load into your model word embeddings
- ▶ Let's look at both.

6.1 Working with text data

LEARNING WORD EMBEDDINGS WITH THE EMBEDDING LAYER

- ▶ choose the vector at **random** - embedding space has **no structure**: the interchangeable words *accurate* and *exact* end up with completely **different embeddings**
- ▶ **synonyms** to be embedded into similar word vectors
- ▶ **geometric distance** (such as L2 distance) between any two word vectors to relate to the **semantic distance** between the associated words



6.1 Working with text data

LEARNING WORD EMBEDDINGS WITH THE EMBEDDING LAYER

- ▶ *cat*, *dog*, *wolf*, and *tiger* - semantic relationships between these words can be encoded as geometric transformations.
- ▶ “from pet to wild animal” - from *cat* to *tiger* and from *dog* to *wolf*
- ▶ “from canine to feline” vector from *dog* to *cat* and from *wolf* to *tiger*
- ▶ “gender” and “plural” vectors - “female” vector + vector “king” → vector “queen,” “plural” vector + vector “king” → “kings.”
- ▶ Word-embedding spaces - interpretable and potentially useful vectors.

Word vectors	Dimensions				
	dog	-0.4	0.37	0.02	-0.34
	cat	-0.15	-0.02	-0.23	-0.23
	lion	0.19	-0.4	0.35	-0.48
	tiger	-0.08	0.31	0.56	0.07
	elephant	-0.04	-0.09	0.11	-0.06
	cheetah	0.27	-0.28	-0.2	-0.43
	monkey	-0.02	-0.67	-0.21	-0.48
	rabbit	-0.04	-0.3	-0.18	-0.47
	mouse	0.09	-0.46	-0.35	-0.24
	rat	0.21	-0.48	-0.56	-0.37




Figure 6.3 A toy example of a word-embedding space

6.1 Working with text data

- ▶ learning the **weights** of a layer: the **Embedding** layer

Listing 6.5 Instantiating an Embedding layer

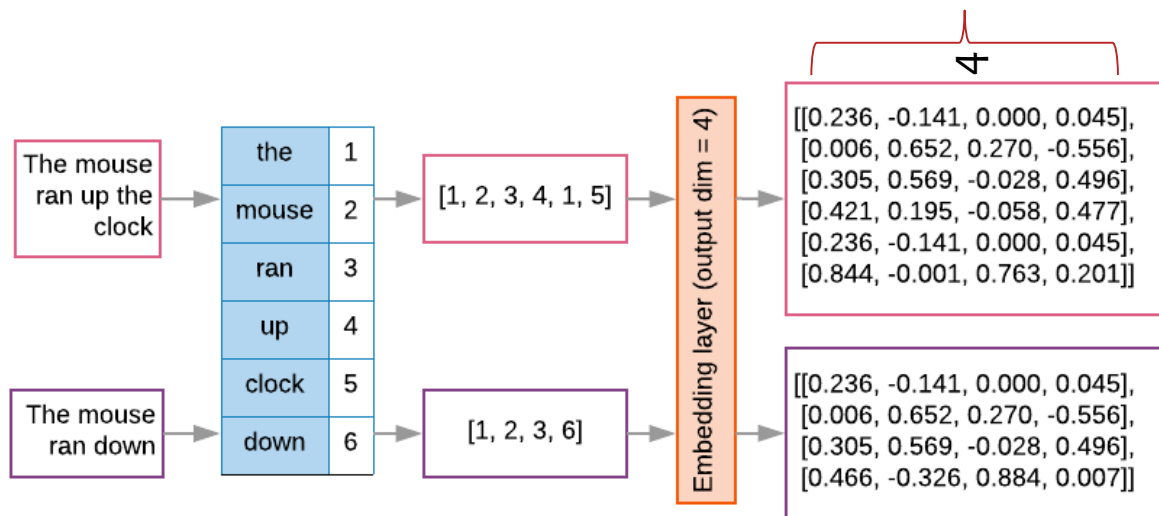
```
from keras.layers import Embedding  
embedding_layer = Embedding(1000, 64)  
#(batch, input_length)
```

- ▶ The **Embedding** layer is best understood as a dictionary that maps **integer indices** (which stand for specific words) to **dense vectors**.
- ▶ It takes integers as **input**, it looks up these integers in an internal **dictionary**, and it returns the associated **vectors**. It's effectively a dictionary lookup.

Word index → **Embedding layer** → Corresponding word **vector**

6.1 Working with text data

- ▶ The Embedding layer takes as input a 2D tensor of integers, of shape (samples, `sequence_length`), where each entry is a sequence of integers.
- ▶ It can embed sequences of variable lengths: (32, 10) (batch of 32 sequences of length 10) or (64, 15) (batch of 64 sequences of length 15).
- ▶ All sequences in a batch must have the **same length**, though (because you need to pack them into a single tensor), so sequences that are shorter than others should be **padded with 0s**, and sequences that are longer should be **truncated**.



6.1 Working with text data

- ▶ This layer **returns** a 3D floating-point tensor of shape (samples, sequence_length, **embedding_dimensionality**).
- ▶ Such a 3D tensor can then be processed by an **RNN layer** or a **1D convolution layer** (both will be introduced in the following sections).
- ▶ Embedding layer - its **weights** (its internal dictionary of token vectors) are initially random → gradually adjusted via backpropagation → embedding space (specialized for the specific problem)
- ▶ IMDB movie-review sentiment-prediction - the **top 10,000** most common words and cut off the reviews after only **20 words**.
- ▶ **input** integer sequences (2D integer tensor) → **embedded sequences** (3D float tensor) → **flatten** the tensor to 2D → **train** a single Dense layer on top for classification → **8-dimensional embeddings** for each of the 10,000 words

6.1 Working with text data

Listing 6.6 Loading the IMDB data for use with an Embedding layer

```
from keras.datasets import imdb
from keras import preprocessing

max_features = 10000 # Number of words to consider as features
maxlen = 20 # Cuts off the text after this number of words
              #(among the max_features most common words)

(x_train, y_train), (x_test, y_test) = imdb.load_data(
    num_words=max_features) # Loads the data as lists of integers

x_train = preprocessing.sequence.pad_sequences(x_train, maxlen=maxlen)
# lists of integers → a 2D integer tensor of shape (samples, maxlen)
x_test = preprocessing.sequence.pad_sequences(x_test, maxlen=maxlen)
# padded with 0s for shorter sequences or truncated for longer sequences
```

6.1 Working with text data

Listing 6.7 Using an Embedding layer and classifier on the IMDB data

```
from keras.models import Sequential
from keras.layers import Flatten, Dense

model = Sequential()

model.add(Embedding(10000, 8, input_length=maxlen))
# Specifies the maximum input length to the Embedding layer
# so you can later flatten the embedded inputs.
# Output of the activations have shape (samples, maxlen, 8) of 3D with 8 Output.

model.add(Flatten()) # 160
# Flattens the 3D tensor of embeddings into a 2D tensor of shape (samples, maxlen * 8)
model.add(Dense(1, activation='sigmoid')) # Adds the classifier on top
model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])
model.summary()

history = model.fit(x_train, y_train,
                    epochs=10, batch_size=32, validation_split=0.2)
# training dataset-8000개, test dataset-2000개
```

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, 20, 8)	80000
flatten_1 (Flatten)	(None, 160)	0
dense_1 (Dense)	(None, 1)	161



6.1 Working with text data

- ▶ You get to a validation accuracy of $\sim 76\%$, which is pretty good considering that you're only looking at the first **20 words** in every review.
- ▶ no **inter-word relationships** and **sentence structure** (for example, this model would likely treat both “**this movie is a bomb**” and “**this movie is the bomb-짱**” as being negative reviews).
- ▶ It's much better to **add recurrent layers** or **1D convolutional layers** on top of the embedded sequences to learn features.


6.1 Working with text data

USING PRETRAINED WORD EMBEDDINGS

- ▶ little training data?
- ▶ precomputed embedding space - highly structured and exhibits useful properties by using word-occurrence statistics, using a variety of techniques, some involving neural networks.
- ▶ Word2vec algorithm (<https://code.google.com/archive/p/word2vec>), developed by Tomas Mikolov at Google in 2013.
 - ▶ Word2vec dimensions capture specific semantic properties, such as genders



(Mikolov et al., NAACL HLT, 2013)



6.1 Working with text data

USING PRETRAINED WORD EMBEDDINGS

► **GloVe**, <https://nlp.stanford.edu/projects/glove>, by Stanford researchers in 2014.

- factorizing a matrix of word **co-occurrence statistics** obtained from millions of English tokens, **Wikipedia** data and **Common Crawl** data.

Window based co-occurrence matrix

- Example corpus:
 - I like deep learning.
 - I like NLP.
 - I enjoy flying.

counts	I	like	enjoy	deep	learning	NLP	flying	.
I	0	2	1	0	0	0	0	0
like	2	0	0	1	0	1	0	0
enjoy	1	0	0	0	0	0	1	0
deep	0	1	0	0	1	0	0	0
learning	0	0	0	1	0	0	0	1
NLP	0	1	0	0	0	0	0	1
flying	0	0	1	0	0	0	0	1
.	0	0	0	0	1	1	1	0

6.1 Working with text data

6.1.3 Putting it all together: from raw text to word embeddings

- ▶ pretrained word embeddings
- ▶ the original text data instead of using the pretokenized IMDB data packaged in Keras

DOWNLOADING THE IMDB DATA AS RAW TEXT

- ▶ download the raw IMDB dataset from <http://mng.bz/0tIo>.
- ▶ Uncompress it.
- ▶ collect the individual training reviews into a list of strings, one string per review.
- ▶ collect the review labels (positive/negative) into a labels list.

6.1 Working with text data

Listing 6.8 Processing the labels of the raw IMDB data

```
import os

imdb_dir = '/Users/fchollet/Downloads/aclImdb'
          # deep-learning-with-python-notebooks-master
train_dir = os.path.join(imdb_dir, 'train')

labels = []
texts = []

for label_type in ['neg', 'pos']: # read data from train_dir: /pos 12,500, /neg 12,500
    dir_name = os.path.join(train_dir, label_type) # ../neg or ../pos
    for fname in os.listdir(dir_name):
        if fname[-4:] == '.txt':
            f = open(os.path.join(dir_name, fname))
            texts.append(f.read())
            f.close()
            if label_type == 'neg':
                labels.append(0)
            else:
                labels.append(1)
```


6.1 Working with text data

TOKENIZING THE DATA

▶ **pretrained word embeddings** - restricting the training data to the first 200 samples (otherwise, task-specific embeddings are likely to outperform)

Listing 6.9 Tokenizing the text of the raw IMDB data

```
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
import numpy as np

maxlen = 100 # Cuts off reviews after 100 words
training_samples = 200 # Trains on 200 samples
validation_samples = 10000 # Validates on 10,000 samples
max_words = 10000 # Considers only the top 10,000 words in the dataset
tokenizer = Tokenizer(num_words=max_words)
tokenizer.fit_on_texts(texts) # 입력에 맞게 내부 list 생성
sequences = tokenizer.texts_to_sequences(texts) # 단어 인덱스만 가져옴
word_index = tokenizer.word_index # 88,582 unique words, 모든 단어 포함
print('Found %s unique tokens.' % len(word_index))

data = pad_sequences(sequences, maxlen=maxlen)
labels = np.asarray(labels)

print('Shape of data tensor:', data.shape)
print('Shape of label tensor:', labels.shape)

Found 88582 unique tokens
Shape of data tensor:(25000,100)
Shape of label tensor :(25000,)
```

```
>>> tokenizer.word_index
{'flour': 3319, 'transit': 2679,
 'in': 16783, 'grips': 4847, 'corru
B, 'md': 32663, 'volleyfoot': 346
```

6.1 Working with text data

```
indices = np.arange(data.shape[0]) # 25,000 - [0 1 2 ... 24997 24998 24999]
# first shuffles the data, all negative first, then all positive
np.random.shuffle(indices) # [23739 2813 974 ... 167 23722 19124]
data = data[indices]
labels = labels[indices]
x_train = data[:training_samples] # 200
y_train = labels[:training_samples] # 200
x_val = data[training_samples:
              training_samples + validation_samples] # 10,000
y_val = labels[training_samples:
                training_samples + validation_samples] # 10,000

x_val: (10000, 100) # maxlen = 100
x_val: [[ 128 1480 413 ... 188 335 543] [ 7 11 6 ... 52 867 97] [ 23 1487 14 ... 2 65 2776]
          ... [ 0 0 0 ... 42 35 615] [ 480 2 327 ... 39 568 3920] [9141 59 1463 ... 128 232 4572]]
y_val: (10000,)
y_val: [0 1 1 ... 1 0 1]
```

6.1 Working with text data

DOWNLOADING THE GLOVE WORD EMBEDDINGS

- ▶ Go to <https://nlp.stanford.edu/projects/glove>, and download the precomputed embeddings from 2014 English Wikipedia. It's an 822 MB zip file called glove.6B.zip, containing **100-dimensional** embedding vectors for **400,000 words** (or nonword tokens). Unzip it.
- ▶ Let's parse the unzipped file (a .txt file) to build an index that maps words (as strings) to their vector representation (as number vectors).

Listing 6.10 Parsing the GloVe word-embeddings file

```
glove_dir='/Users/fchollet/Downloads/glove.6B' #word+100 dim vector
embeddings_index = {} # 400,000 words
f = open(os.path.join(glove_dir, 'glove.6B.100d.txt'))
for line in f:
    values = line.split()
    word = values[0]
    coefs = np.asarray(values[1:], dtype='float32')
    embeddings_index[word] = coefs
f.close()

print('Found %s word vectors.' % len(embeddings_index))
# Found 400000 word vectors
print(len(embeddings_index.get('cat')))#100-dimensional embedding vectors
# 100
# [0.23088 0.28283 ... -0.71493 ]
# Word = sandberger
# Coefs = [0.28365 -0.6263,,, -0.15701]
```

6.1 Working with text data

- ▶ Next, you'll build an embedding matrix that you can load into an **Embedding layer**. It must be a matrix of shape `(max_words, embedding_dim)` - `(10000, 100)`
- ▶ Note that **index 0** isn't supposed to stand for any **word or token**—it's a placeholder.

```
# embeddings_index = {} # 400,000 words
```

Listing 6.11 Preparing the GloVe word-embeddings matrix

```
embedding_dim = 100 # len(embeddings_index.get('cat'))

embedding_matrix = np.zeros((max_words, embedding_dim)) # max_words=10000
for word, i in word_index.items(): # [('the', 1), ('and', 2), ..., ('hued', 88582)]
    if i < max_words: # max_words=10000
        embedding_vector = embeddings_index.get(word) #
        if embedding_vector is not None: # exist in embeddings_index
            embedding_matrix[i] = embedding_vector
# else - Words not found in the embedding index will be all zeros.
# ... [0.0.0. ... 0.0.][[-0.038194 -0.24487001 ... 0.27061999] - (10000, 100)
```

6.1 Working with text data

DEFINING A MODEL

- ▶ Let's parse the unzipped file (a .txt file) to build an index that maps words (as strings) to their vector representation (as number vectors).

Listing 6.12 Model definition

```
from keras.models import Sequential
from keras.layers import Embedding, Flatten, Dense

model = Sequential()
model.add(Embedding(max_words, embedding_dim, #max_words=10000, embedding_dim=100
                    input_length=maxlen)) # maxlen = 100, p.25: pad_sequences
model.add(Flatten())
model.add(Dense(32, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
model.summary()
```

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, 100, 100)	1000000
flatten_1 (Flatten)	(None, 10000)	0
dense_1 (Dense)	(None, 32)	320032
dense_2 (Dense)	(None, 1)	33

6.1 Working with text data

LOADING THE GLOVE EMBEDDINGS IN THE MODEL

- ▶ The **Embedding layer** has a single weight matrix: a 2D float matrix where each entry i is the word vector meant to be associated with index i .
- ▶ Load the **GloVe matrix** you prepared into the **Embedding layer**, the first layer in the model
- ▶ **freeze the Embedding layer** (set its trainable attribute to **False**)

Listing 6.13 Loading pretrained word embeddings

```
model.layers[0].set_weights([embedding_matrix])  
#                               (10000, 100)  
model.layers[0].trainable = False
```

6.1 Working with text data

TRAINING AND EVALUATING THE MODEL

- ▶ Compile and train the model.

Listing 6.14 Training and evaluation into the Embedding

```
model.compile(optimizer='rmsprop',  
              loss='binary_crossentropy', metrics=['acc'])  
history = model.fit(x_train, y_train,  
                    epochs=10, batch_size=32,  
                    validation_data=(x_val, y_val))  
model.save_weights('pre_trained_glove_model.h5')
```

6.1 Working with text data

- Now, plot the model's performance over time (see figures 6.5 and 6.6)

Listing 6.15 Plotting the results layer

```
import matplotlib.pyplot as plt
acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy') plt.legend()

plt.figure()

plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()

plt.show()
```

Figure 6.5 Training and validation loss when using pretrained word embeddings





Figure 6.6 Training and validation accuracy when using pretrained word embeddings

6.1 Working with text data

- ▶ The model quickly starts overfitting, which is unsurprising given the small number of training samples. Validation accuracy has high variance for the same reason, but it seems to reach the **high 50s**.
- ▶ so few training samples - performance is heavily dependent on exactly which **200 samples** you choose randomly.
- ▶ **without loading the pretrained word embeddings** and without freezing the embedding layer. In that case, you'll learn a task-specific embedding of the input tokens, which is generally more powerful than pretrained word embeddings when lots of data is available.
- ▶ But in this case, you have only **200 training samples**. Let's try it (see figures 6.7 and 6.8).

6.1 Working with text data

Listing 6.16 Training the same model *without pretrained word embeddings*

```
from keras.models import Sequential
from keras.layers import Embedding, Flatten, Dense

model = Sequential()
model.add(Embedding(max_words, embedding_dim, input_length=maxlen))
model.add(Flatten())
model.add(Dense(32, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
model.summary()

model.compile(optimizer='rmsprop',
              loss='binary_crossentropy', metrics=['acc'])
history = model.fit(x_train, y_train,
                    epochs=10, batch_size=32,
                    validation_data=(x_val, y_val))
```





Figure 6.7 Training and validation loss without using pretrained word embeddings

Figure 6.8 Training and validation accuracy without using pretrained word embeddings



6.1 Working with text data

- ▶ Validation accuracy stalls in the **low 50s**. So in this case, **pretrained word embeddings** outperform jointly learned embeddings. If you **increase the number of training samples**, this will quickly stop being the case—try it as an exercise.
- ▶ Finally, let's evaluate the model on the **test data**. First, you need to tokenize the test data.

Listing 6.17 Tokenizing the data of the test set

```
test_dir = os.path.join(imdb_dir, 'test')
labels = []
texts = []

for label_type in ['neg', 'pos']:
    dir_name = os.path.join(test_dir, label_type)
    for fname in sorted(os.listdir(dir_name)):
        if fname[-4:] == '.txt':
            f = open(os.path.join(dir_name, fname))
            texts.append(f.read())
            f.close()
            if label_type == 'neg':
                labels.append(0)
            else:
                labels.append(1)

sequences = tokenizer.texts_to_sequences(texts)
x_test = pad_sequences(sequences, maxlen=maxlen)
y_test = np.asarray(labels)
```

6.1 Working with text data

- ▶ Next, load and evaluate the first model

Listing 6.18 Evaluating the model on the test set

```
model.load_weights('pre_trained_glove_model.h5')  
model.evaluate(x_test, y_test)
```


- ▶ You get an appalling test accuracy of **56%**. Working with just a handful of training samples is **difficult**!
 - Turn raw text into something a neural network can process
 - Use the **Embedding layer** in a Keras model to learn task-specific token embeddings
 - Use **pretrained word embeddings** to get an extra boost on small natural-language-processing problems
- ▶ 실습 - raw IMDB data를 GloVe word-embeddings file을 이용하여 다음과 같이 변경하여 분류하세요
 - ▶ 변경 가능한 변수 조정
 - ▶ Train data 조정
 - ▶ 모델 수정
 - ▶ 결과 - 분류율과 분석 결과

6.2 Understanding recurrent neural networks

- ▶ A **major characteristic** of all neural networks is that they have **no memory** - *no state kept in **between inputs***
- ▶ **feedforward networks** - IMDB example: **an entire movie review** was transformed into **a single large vector** and processed in one go.
- ▶ **word by word** (eye saccade by eye saccade) - from **past information** → constantly updated **as new information**

6.2 Understanding recurrent neural networks

► A *recurrent neural network* (RNN) - it processes sequences by *iterating* through the sequence elements and maintaining a *state* containing information relative to what it has seen so far.

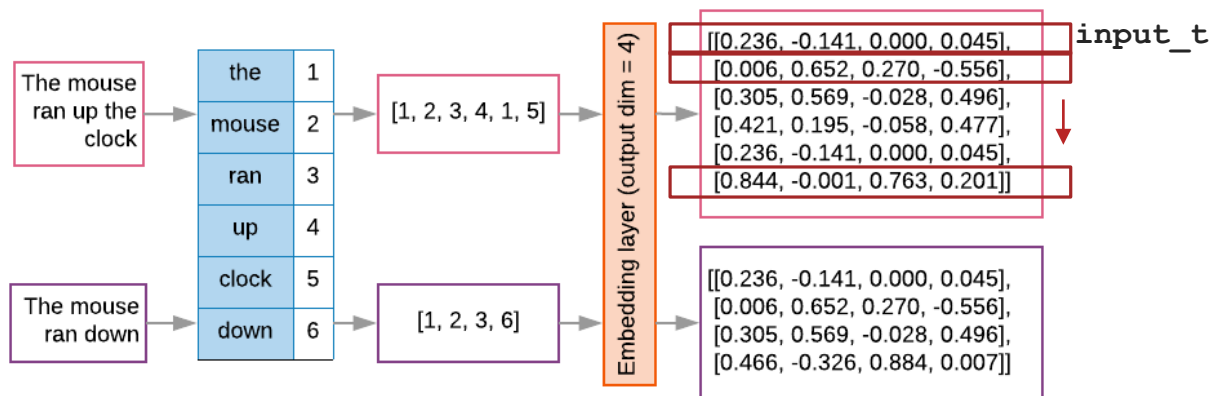


6.2 Understanding recurrent neural networks

- ▶ RNN takes as input a sequence of vectors - 2D tensor of size (timesteps, input_features)
- ▶ set the **state** for the next step to be this previous output.
- ▶ *initial state* - **all-zero vector**

Listing 6.19 Pseudocode RNN

```
state_t = 0 # The state at t
for input_t in input_sequence:
    # Iterates over sequence elements
    output_t = f(input_t, state_t)
    state_t = output_t
    # The previous output becomes
    # the state for the next iteration.
```




6.2 Understanding recurrent neural networks

► f : input and state $\rightarrow W$ and U , and a bias vector

Listing 6.20 More detailed pseudocode for the RNN

```
state_t = 0 # h
for input_t in input_sequence:
    output_t = activation(dot(W, input_t) + dot(U, state_t) + b)
    state_t = output_t
```



6.2 Understanding recurrent neural networks


► naive Numpy implementation of the forward pass of the simple RNN.

Listing 6.21 Numpy implementation of a simple RNN

```
import numpy as np

timesteps = 100 # Number of timesteps in the input sequence
input_features = 32 # Dimensionality of the input feature space
output_features = 64 # Dimensionality of the output feature space
inputs = np.random.random((timesteps, input_features))
# Input data: random noise for the sake of the example
state_t = np.zeros((output_features,))
# Initial state: all-0 vector

W = np.random.random((output_features, input_features)) # (64,32)
U = np.random.random((output_features, output_features)) # (64,64)
b = np.random.random((output_features,)) # (64,)
# 1 random weight matrices
```



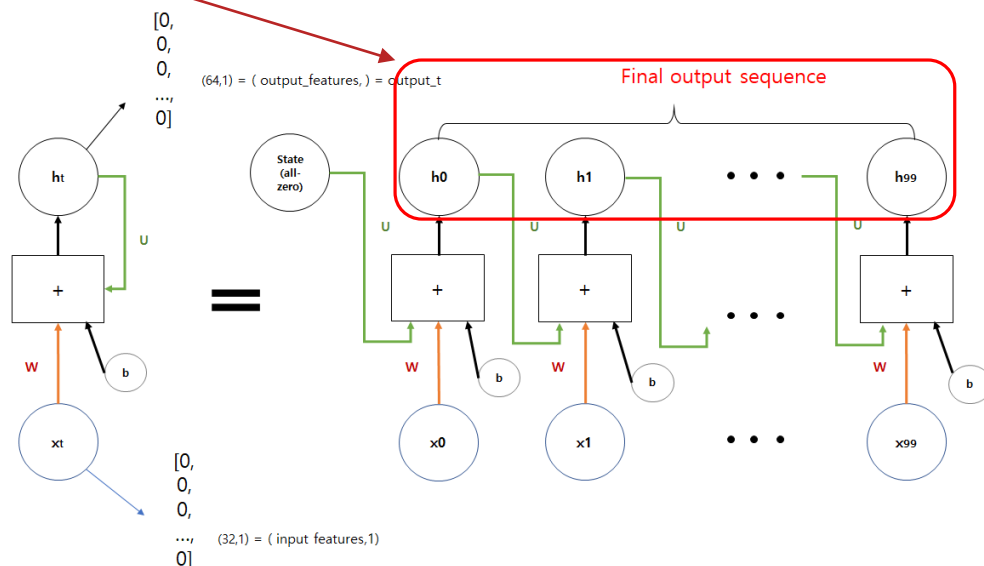

6.2 Understanding recurrent neural networks

```

successive_outputs = []
for input_t in inputs: # a vector of shape (input_features,)
    output_t = np.tanh(np.dot(W, input_t) + np.dot(U, state_t) + b)
    # Combines the input with the current state (the previous output)
    successive_outputs.append(output_t) #Stores this output in a list
    state_t = output_t
    # Updates the state of the network for the next timestep
final_output_sequence = np.concatenate(successive_outputs, axis=0)
# The final output is a 2D tensor of
# shape (timesteps, output_features), shape = (100, 64)

```

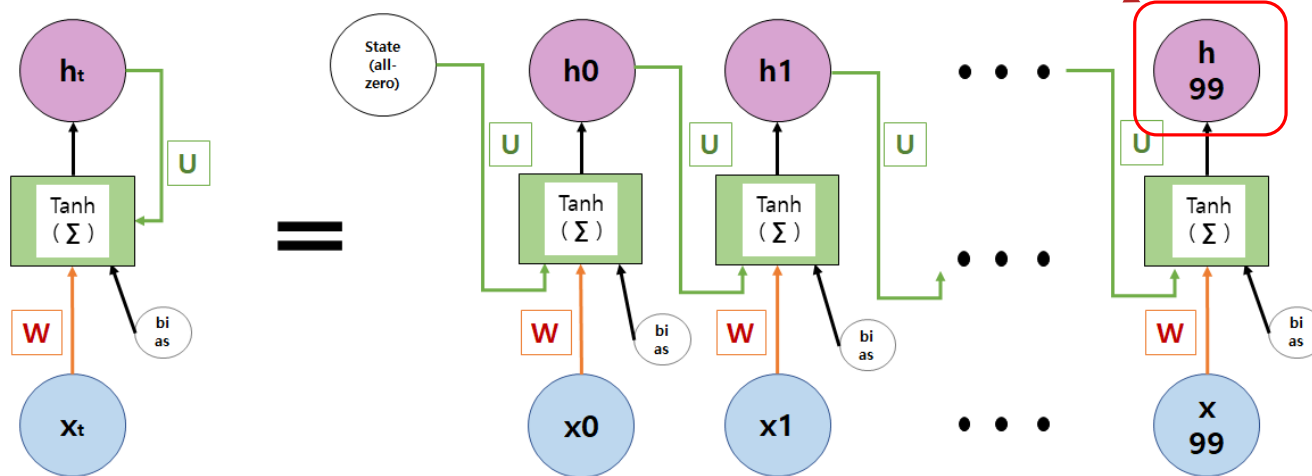
Setting axis=0 concatenates along the row axis



6.2 Understanding recurrent neural networks

`output_t = np.tanh(np.dot(W, input_t) + np.dot(U, state_t) + b)`

NOTE The final output is a 2D tensor of shape (timesteps, output_features) at time t . Only the last output (`output_t`) at the end of the loop is needed, because it already contains information about the entire sequence.



6.2 Understanding recurrent neural networks

6.2.1 A recurrent layer in Keras

- ▶ SimpleRNN layer:

```
from keras.layers import SimpleRNN
```

- ▶ There is one minor difference: SimpleRNN processes **batches of sequences**, like all other Keras layers

(timesteps, input_features) →

(batch_size, timesteps, input_features)

- ▶ two different modes of **return**

- ▶ (batch_size, timesteps, **output_features**) - the full sequences of successive outputs
- ▶ (batch_size, **output_features**) - only the **last output** for each input sequence

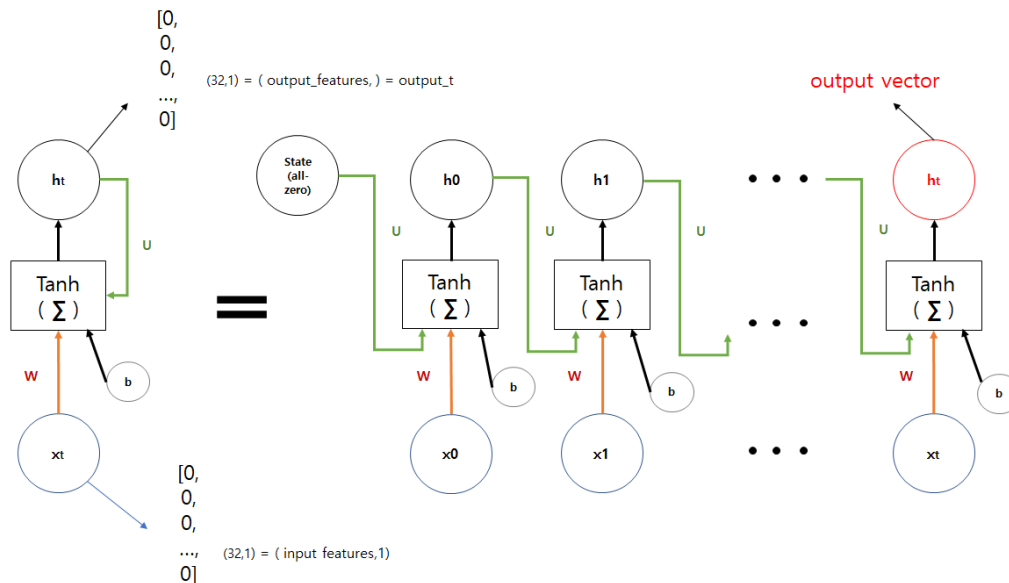
- ▶ These two modes are controlled by the **return_sequences** constructor argument.

6.2 Understanding recurrent neural networks

- SimpleRNN and returns only the output at the **last timestep**:

```
>>> from keras.models import Sequential
>>> from keras.layers import Embedding, SimpleRNN
>>> model = Sequential()
>>> model.add(Embedding(10000, 32)) # (max_features, output dim)
>>> model.add(SimpleRNN(32))
>>> model.summary()
```

Layer (type)	Output Shape	Param #
embedding_22 (Embedding)	(None, None, 32)	320000
simplernn_10 (SimpleRNN)	(None, 32)	2080 = (32*32*2+32)

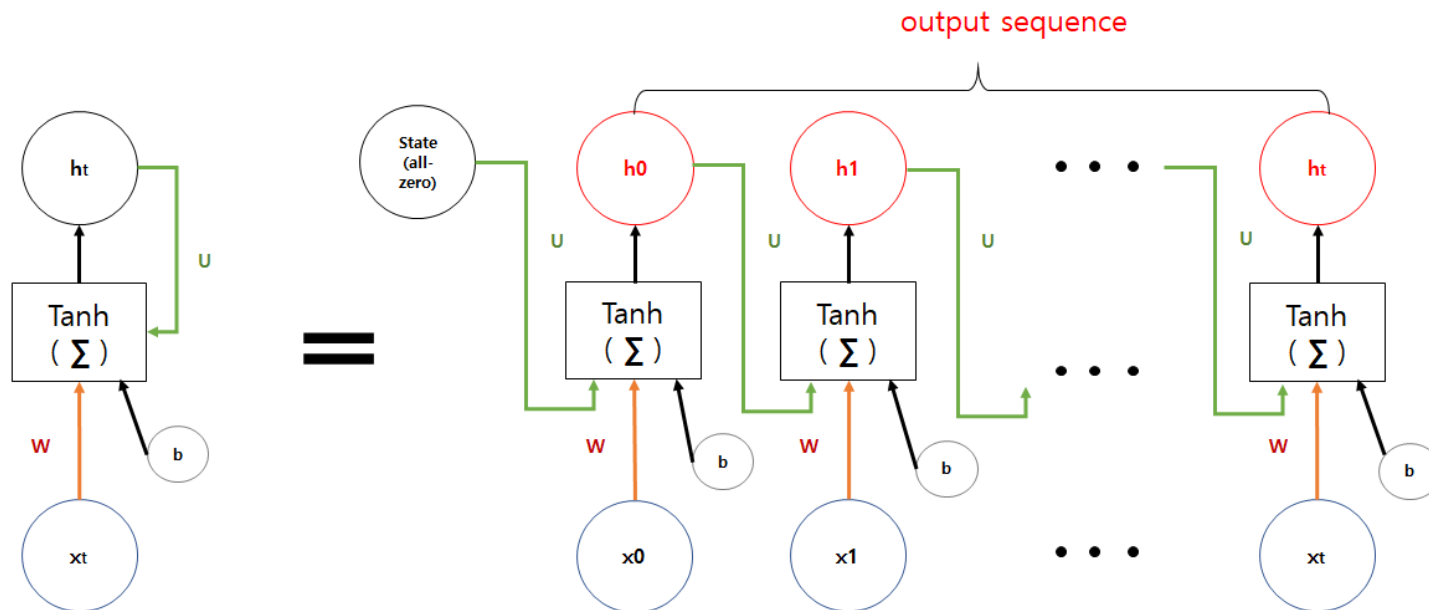


6.2 Understanding recurrent neural networks

- ▶ The following example returns the **full state sequence**:

```
>>> model = Sequential()  
>>> model.add(Embedding(10000, 32))  
>>> model.add(SimpleRNN(32, return_sequences=True))  
>>> model.summary()
```

Layer (type)	Output Shape	Param #
embedding_23 (Embedding)	(None, None, 32)	320000
simplernn_11 (SimpleRNN)	(None, None, 32)	2080



6.2 Understanding recurrent neural networks

- ▶ stack several recurrent layers:

```
>>> model = Sequential()
>>> model.add(Embedding(10000, 32))
>>> model.add(SimpleRNN(32, return_sequences=True))
>>> model.add(SimpleRNN(32, return_sequences=True))
>>> model.add(SimpleRNN(32, return_sequences=True))
>>> model.add(SimpleRNN(32)) # Last layer only returns the last output
>>> model.summary()
```

Layer (type)	Output Shape	Param #
embedding_3 (Embedding)	(None, None, 32)	320000
simple_rnn_3 (SimpleRNN)	(None, None, 32)	
simple_rnn_4 (SimpleRNN)	(None, None, 32)	2080
simple_rnn_5 (SimpleRNN)	(None, None, 32)	2080
simple_rnn_6 (SimpleRNN)	(None, 32)	2080

6.2 Understanding recurrent neural networks

- ▶ IMDB movie-review-classification problem - First, preprocess the data.

Listing 6.22 Preparing the IMDB data

```
from keras.datasets import imdb
from keras.preprocessing import sequence
max_features = 10000 # Number of words to consider as features
maxlen = 500 # Cuts off texts after this many words
batch_size = 32
(input_train, y_train), (input_test, y_test) =
    imdb.load_data(num_words=max_features)
print(len(input_train), 'train sequences') # 25000
print(len(input_test), 'test sequences') # 25000
input_train = sequence.pad_sequences(input_train, maxlen=maxlen)
input_test = sequence.pad_sequences(input_test, maxlen=maxlen)
print('input_train shape:', input_train.shape) # (25000, 500)
print('input_test shape:', input_test.shape) # (25000, 500)
```


6.2 *Understanding recurrent neural networks*

- ▶ Train an Embedding layer and a SimpleRNN layer.

Listing 6.23 Training the model with Embedding and SimpleRNN layers

```
from keras.layers import Dense
model = Sequential()
model.add(Embedding(max_features, 32))
model.add(SimpleRNN(32))
model.add(Dense(1, activation='sigmoid'))
model.compile(optimizer='rmsprop', loss='binary_crossentropy',
              metrics=['acc'])
history = model.fit(input_train, y_train,
                    epochs=10, batch_size=128, validation_split=0.2) #20%
```



6.2 Understanding recurrent neural networks

Listing 6.24 Plotting results

```
import matplotlib.pyplot as plt
acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()

plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```



6.2 Understanding recurrent neural networks

- ▶ In chapter 3, test accuracy - 88%
- ▶ recurrent network - 85% validation accuracy
- ▶ Inputs only the first 500 words - less information than the earlier baseline model.
- ▶ SimpleRNN - No good at processing long sequences, such as text (vanishing information).
- ▶ Other types of recurrent layers perform much better.

6.2 Understanding recurrent neural networks


6.2.2 A Understanding the LSTM and GRU layers

- ▶ LSTM and GRU - SimpleRNN has a major issue: **long-term dependencies** are impossible to learn.
- ▶ This is due to the *vanishing gradient problem*, an effect that is similar to what is observed with non-recurrent networks (feedforward networks) studied by Hochreiter, Schmidhuber, and Bengio in the early 1990s.
- ▶ Long Short-Term Memory (LSTM) algorithm was developed by Hochreiter and Schmidhuber in 1997.
- ▶ **Carry Track** (C) information across many timesteps to save information for later, thus preventing older signals from gradually vanishing during processing.


LSTM-GRU Architecture - overview

<https://towardsdatascience.com/illustrated-guide-to-lstms-and-gru-s-a-step-by-step-explanation-44e9eb85bf21>

LSTM



GRU



sigmoid



tanh



pointwise
multiplication

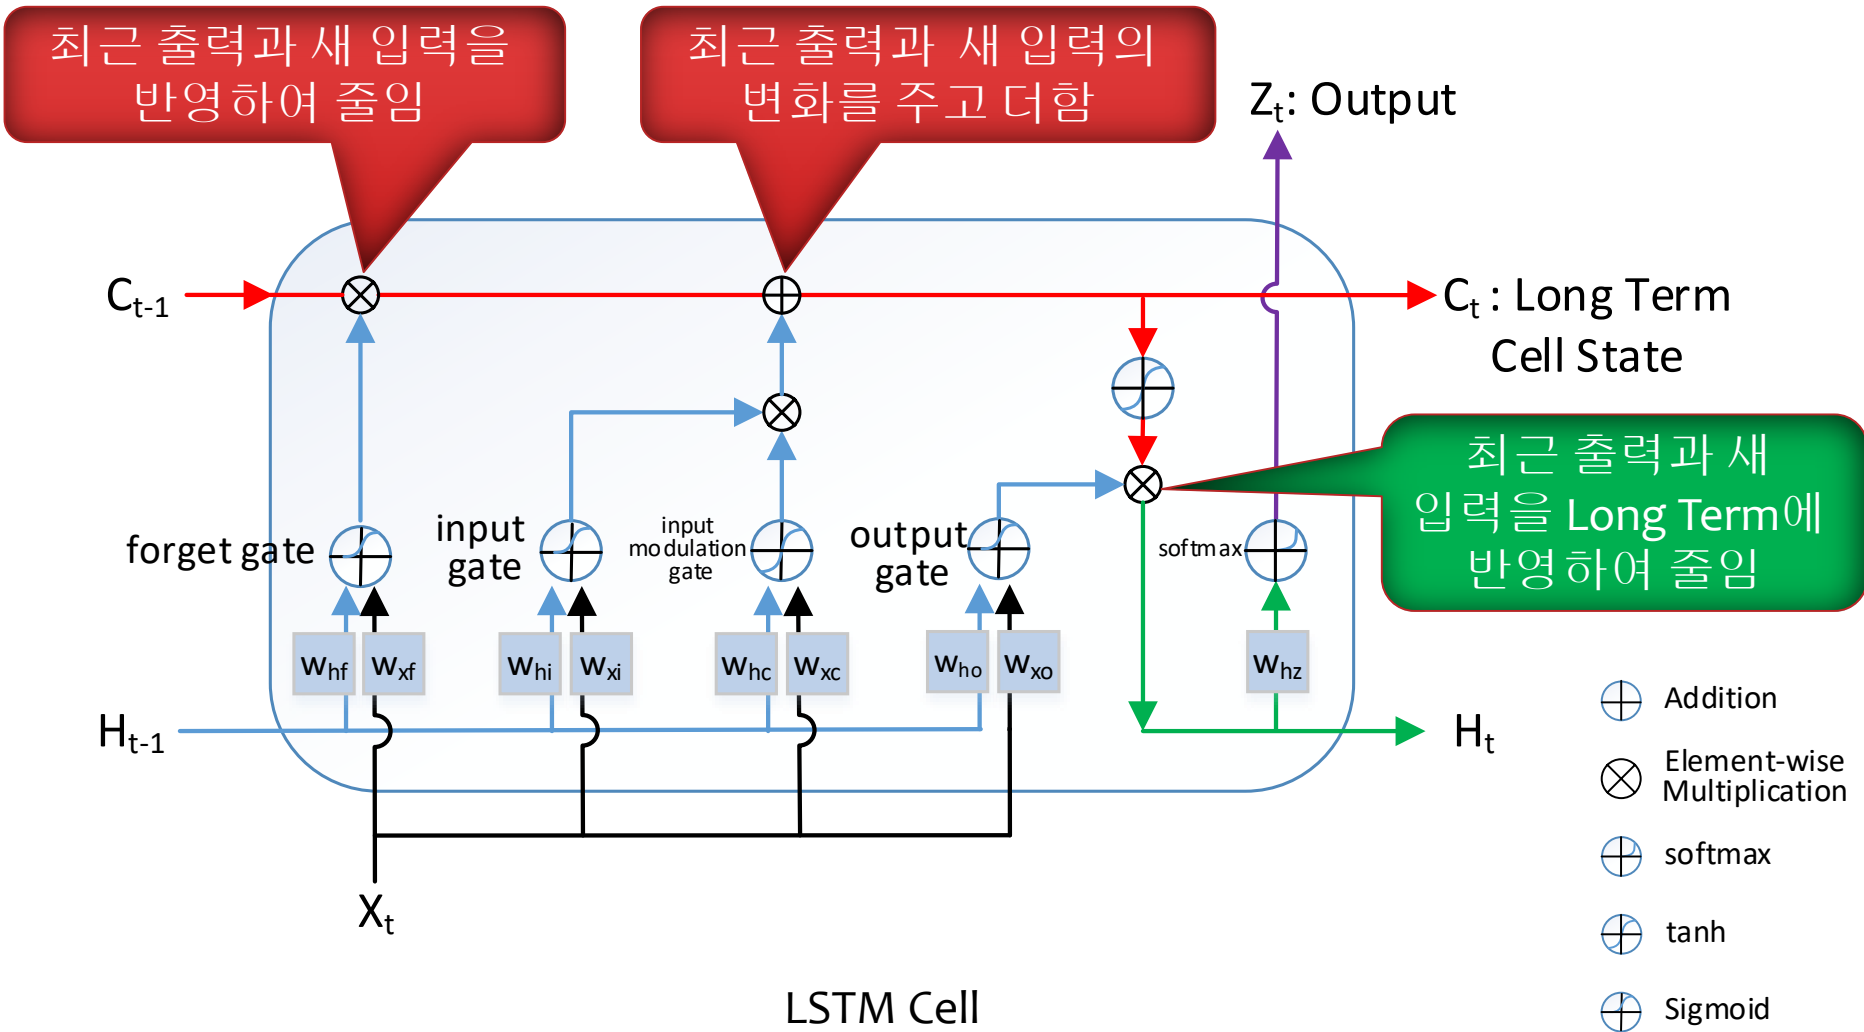


pointwise
addition




vector
concatenation

○ ○ ○ LSTM Architecture - overview ○ ○ ○




○ ○ ○ LSTM Architecture - overview ○ ○ ○




Forget gate operations

○ ○ ○ LSTM Architecture - overview ○ ○ ○




Input gate operations

LSTM Architecture - overview



Calculating cell state

○ ○ ○ LSTM Architecture - overview ○ ○ ○



output gate operations



6.2 Understanding recurrent neural networks

6.2.3 A concrete LSTM example in Keras

- ▶ set up a model using an LSTM layer and train it on the IMDB data (see figures 6.16 and 6.17).
- ▶ similar to the one with SimpleRNN - specify the output dimensionality of the LSTM layer; leave every other argument (there are many) at the Keras defaults.

Listing 6.27 Using the LSTM layer in Keras

```
from keras.layers import LSTM
model = Sequential()
model.add(Embedding(max_features, 32))
model.add(LSTM(32))
model.add(Dense(1, activation='sigmoid'))
model.compile(optimizer='rmsprop',
              loss='binary_crossentropy', metrics=['acc'])
history = model.fit(input_train, y_train,
                    epochs=10, batch_size=128, validation_split=0.2)
```



6.2 Understanding recurrent neural networks

6.2.3 A concrete LSTM example in Keras

- ▶ achieve up to **89%** validation accuracy with less vanishing-gradient problem—and slightly better than the fully connected approach from chapter 3
- ▶ **less data** than you were in chapter 3 by truncating sequences after 500 timesteps, whereas in chapter 3 (10,000), you were considering full sequences.
- ▶ Why isn't LSTM performing better?
 - ▶ no effort to **tune** hyperparameters such as the **embeddings** dimensionality or the **LSTM output** dimensionality.
 - ▶ lack of **regularization**
 - ▶ analyzing the **global, long-term structure of the reviews** (what LSTM is good at) isn't helpful for a **sentiment-analysis** problem.
 - ▶ well solved by looking at what **words occur** in each review, and at what **frequency** in FCN
 - ▶ the strength of LSTM will become apparent: in particular, **question-answering** and **machine translation**

6.2.4 Wrapping up

- ▶ Now you understand the following:
 - What RNNs are and how they work
 - What LSTM is, and why it works better on long sequences than a naive RNN
 - How to use Keras RNN layers to process sequence data
- ▶ Next, advanced features of RNNs